

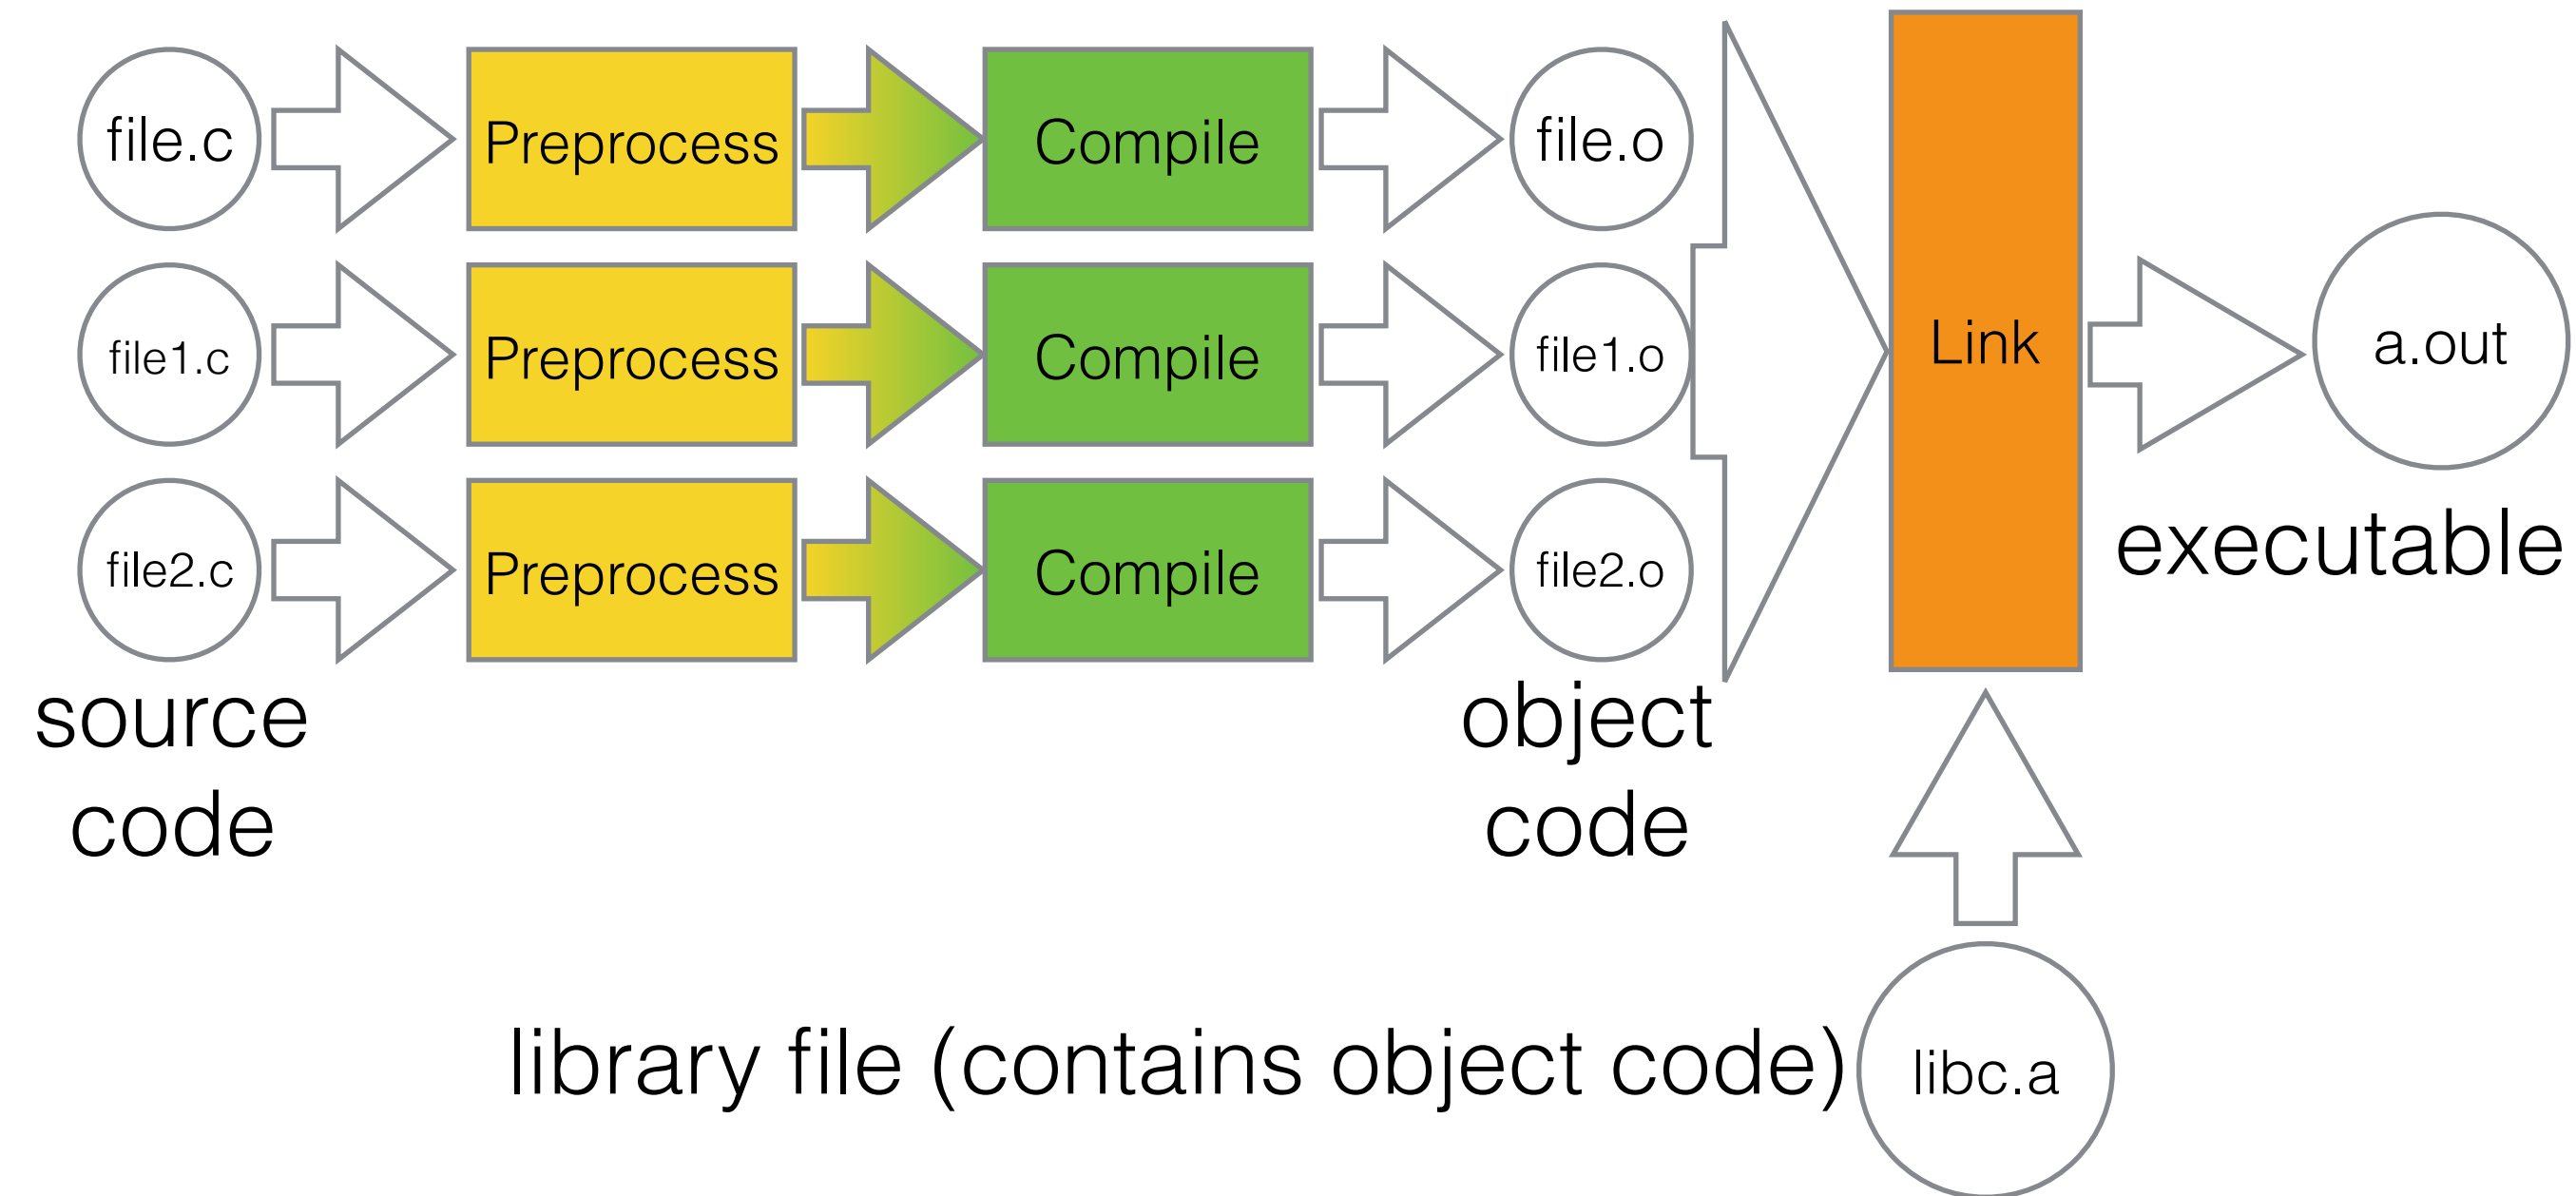
The C Build Process

(P2T: C Lecture 6)

The C Build Process

gcc file.c file1.c file2.c

gcc file.o file1.o file2.o



The Preprocessor

- The preprocessor “prepares” the initial source code for the compiler.
- The output is *also* C source code.
- The preprocessor performs text manipulation on the input source
 - replace text with other text
 - insert or remove text (based on a test)
 - include text from another file at this point
- All *preprocessor directives* start with a #

#define

- `#define` is the simplest preprocessor directive.
- It takes a single word “`macro`”, and replaces it with the `provided text`, everywhere in the rest of the file.
- (Text inside `" "` is not replaced.)

```
#define APPLE DELICIOUS FRUIT
```

```
This is an APPLE, not a SAPPLE.  
"Quoted APPLE text."
```

Doing this with numeric literals is one way to set global constants, like π .

```
#define PI 3.14159
```

preprocess



```
This is an DELICIOUS FRUIT, not a SAPPLE.  
"Quoted APPLE text."
```

"Predefined" macros

- C language provides some macros which will automatically be replaced with useful values by the preprocessor.
- `__LINE__` will be replaced with the line number it appears on.
- `__FILE__` is the name of the file being processed.
- `__DATE__` and `__TIME__` are replaced appropriately.

```
printf("We are on line %d in file %s.\n", __LINE__, __FILE__);
```



```
printf("We are on line %d in file %s.\n", 23, "myfile.c");
```

#ifdef ... #else ... #endif

- User defined macros can also be used to control if a piece of text is kept in, or removed from, the source code, using `#ifdef` constructs.

```
#define APPLE  
  
#ifdef APPLE  
We know about apple.  
#else  
We don't know about apple.  
#endif
```

**APPLE is now #defined,
albeit to an empty value.**

Or

```
#define APPLE  
  
#ifndef APPLE  
We don't know about apple.  
#else  
We know about apple.  
#endif
```

preprocess

We know about apple.

**If APPLE wasn't #defined,
what would the resulting text be?**

preprocess

#include


- **#include** copies all the text from the named file, and inserts it into the text at the point of the directive. (*The preprocessor will then process the included text, if it also contains directives.*)

```
#include "mytext"

This is also some BANANA
```

mytext

```
This is the text in mytext.
It is not very interesting.
#define BANANA text.
```



preprocess

```
This is the text in mytext.
It's not very interesting.

This is also some text.
```

**#define in mytext makes
this change here**

#include <> versus ""

- If you specify the filename in an `#include` like `<stdio.h>` then the preprocessor will look for the file in some standard system locations.
- If you want to refer to a file that's in the same directory as your file (for example), then use `"myfile"`. (You can give a path if you want.)

```
#include <stdio.h>
```

```
#include "myheader.h"
```

**Looks for `stdio.h` file in system locations
for important header files.
(`/usr/include/` etc)**

**Looks for `myheader.h` file in local directory path
(and other "local" locations).**

Common definitions

- In order to allow code in one (source) file to know about functions and values in other (source) files, we need a way to add:
 - function prototypes for functions defined in other files
 - common constants etc
- We can then use those functions in our other files (and let the linker join up all the code at the end).

Header files

- C calls a file which exists just to be `#included` to provide common declarations a "header file".
- Conventionally, they have the `.h` suffix (but they contain normal C code).
- For every `.c` file you write which contains code you want to reference in other `.c` files:
 - Make a `.h` file for the prototypes etc you want to reference elsewhere.

#include guards

- It's an error in C to declare a function (or any other name) more than once in the same file.
- If our definition is in a header, how do we stop the header being included more than once?

file.h

```
#ifndef FILE_H
#define FILE_H

(text of header)

#endif
```

The first time this is #included, the #define activates, and defines FILE_H.

The *second* time we include it, the #ifndef skips the entire file!

**Modern compilers also support putting
#pragma once
at the top of a header file, for same effect.
(This is *not* universally supported!)**

Preprocessing example

example.c

```
#include <stdio.h>
#include "function.h"
```

What happens if we don't include
function.h?
(What does the warning mean?)

```
int main(void){
    printf("The result of function f is: %f\n", f(27));
    return 0;
}
```

function.h

```
//declaration of f (multiplies number by 2)
double f(double);
```

function.c

```
double f(double a) {
    return a*=2;
}
```

```
gcc -E example.c | less
```

(what is the prototype for printf?)

The Compiler

- The compiler takes C source code and translates it into actual machine code.
- Source code is broken down into syntactic units (“tokens”), and the relationships between them.
- Instructions are then converted to machine code, with “symbols” attached to named items for reference.
- The resulting machine code, annotated with symbols, and prepended with a list of all the symbols in the file, is called “object code”.

Optimisation

- As with natural languages, naïve compilation (translating “word for word” into machine code) doesn’t always give the best result.
- We can ask the compiler to spend more time and effort to produce more efficient, or faster, or shorter, representations - *optimisations*.
- We specify these with the `-O` flag, specifying a value from `0` (no optimisations) through `3` (all optimisations).
- In general, `-O2` is a good balance between compilation time and code performance.
- In all cases, the logic (within the C Standard definitions) is preserved - the implementation may just be changed.

Compilation example

example.c

```
#include <stdio.h>
#include "function.h"

int main(void){
    printf("The result of function f is: %d\n", f(27));
    return 0;
}
```

function.h

```
//declaration of f (multiplies number by 2)
int f(int);
```

function.c

```
int f(int a) {
    return a*=2;
}
```

```
gcc -c example.c
gcc -c function.c
```

```
nm example.o
nm function.o
```

The Linker

- The Linker is responsible for joining up the object code from the Compiler, so that all symbols are mapped to their representation.
- Firstly, it takes all of the object code files it is given (one for each C source file), and turns them into one large file, with “consolidated” symbol index at the top.
- Before object code is linked, function calls are simply a note of the symbol that corresponds to the function they want to call.
- The linker looks up the symbol in the symbol index and replaces it with a call to the function code with the corresponding symbol.
- This also happens for variable names in the file scope in each file.

Libraries

- One way to provide a set of useful functions to others is to package them all up into a "library".
- This is just a specially packaged set of object code (with an index to make it easy to find things in it).
- Libraries have names ending in **.a** or **.so** (the difference is in how and when they are linked).
- We can ask the linker to link to a library called *libname.a* (or *libname.so*) with the option **-lname**
- In this case, the linker will also match symbols in our code with symbols in the library mentioned.

The C Standard Library

- One library is used in (almost) all C code you will ever write: `libc.a` (or `libc.so`).
- The "C Standard Library" contains useful functions for a large number of potential applications, including I/O etc.
- It's so large that the prototypes for its functions are split into multiple header files, grouped by topic.
 - `stdio.h` provides prototypes for the I/O functions in `libc`
 - `string.h` provides prototypes for the string functions in `libc`
- As `libc` is so commonly used, the linker will link it without even being asked.

An exception is the (floating point) mathematics part of `libc`. For historical reasons, this is often stored in a separate library `libm.a` / `libm.so` , which *does* need explicitly linked.

Library locations

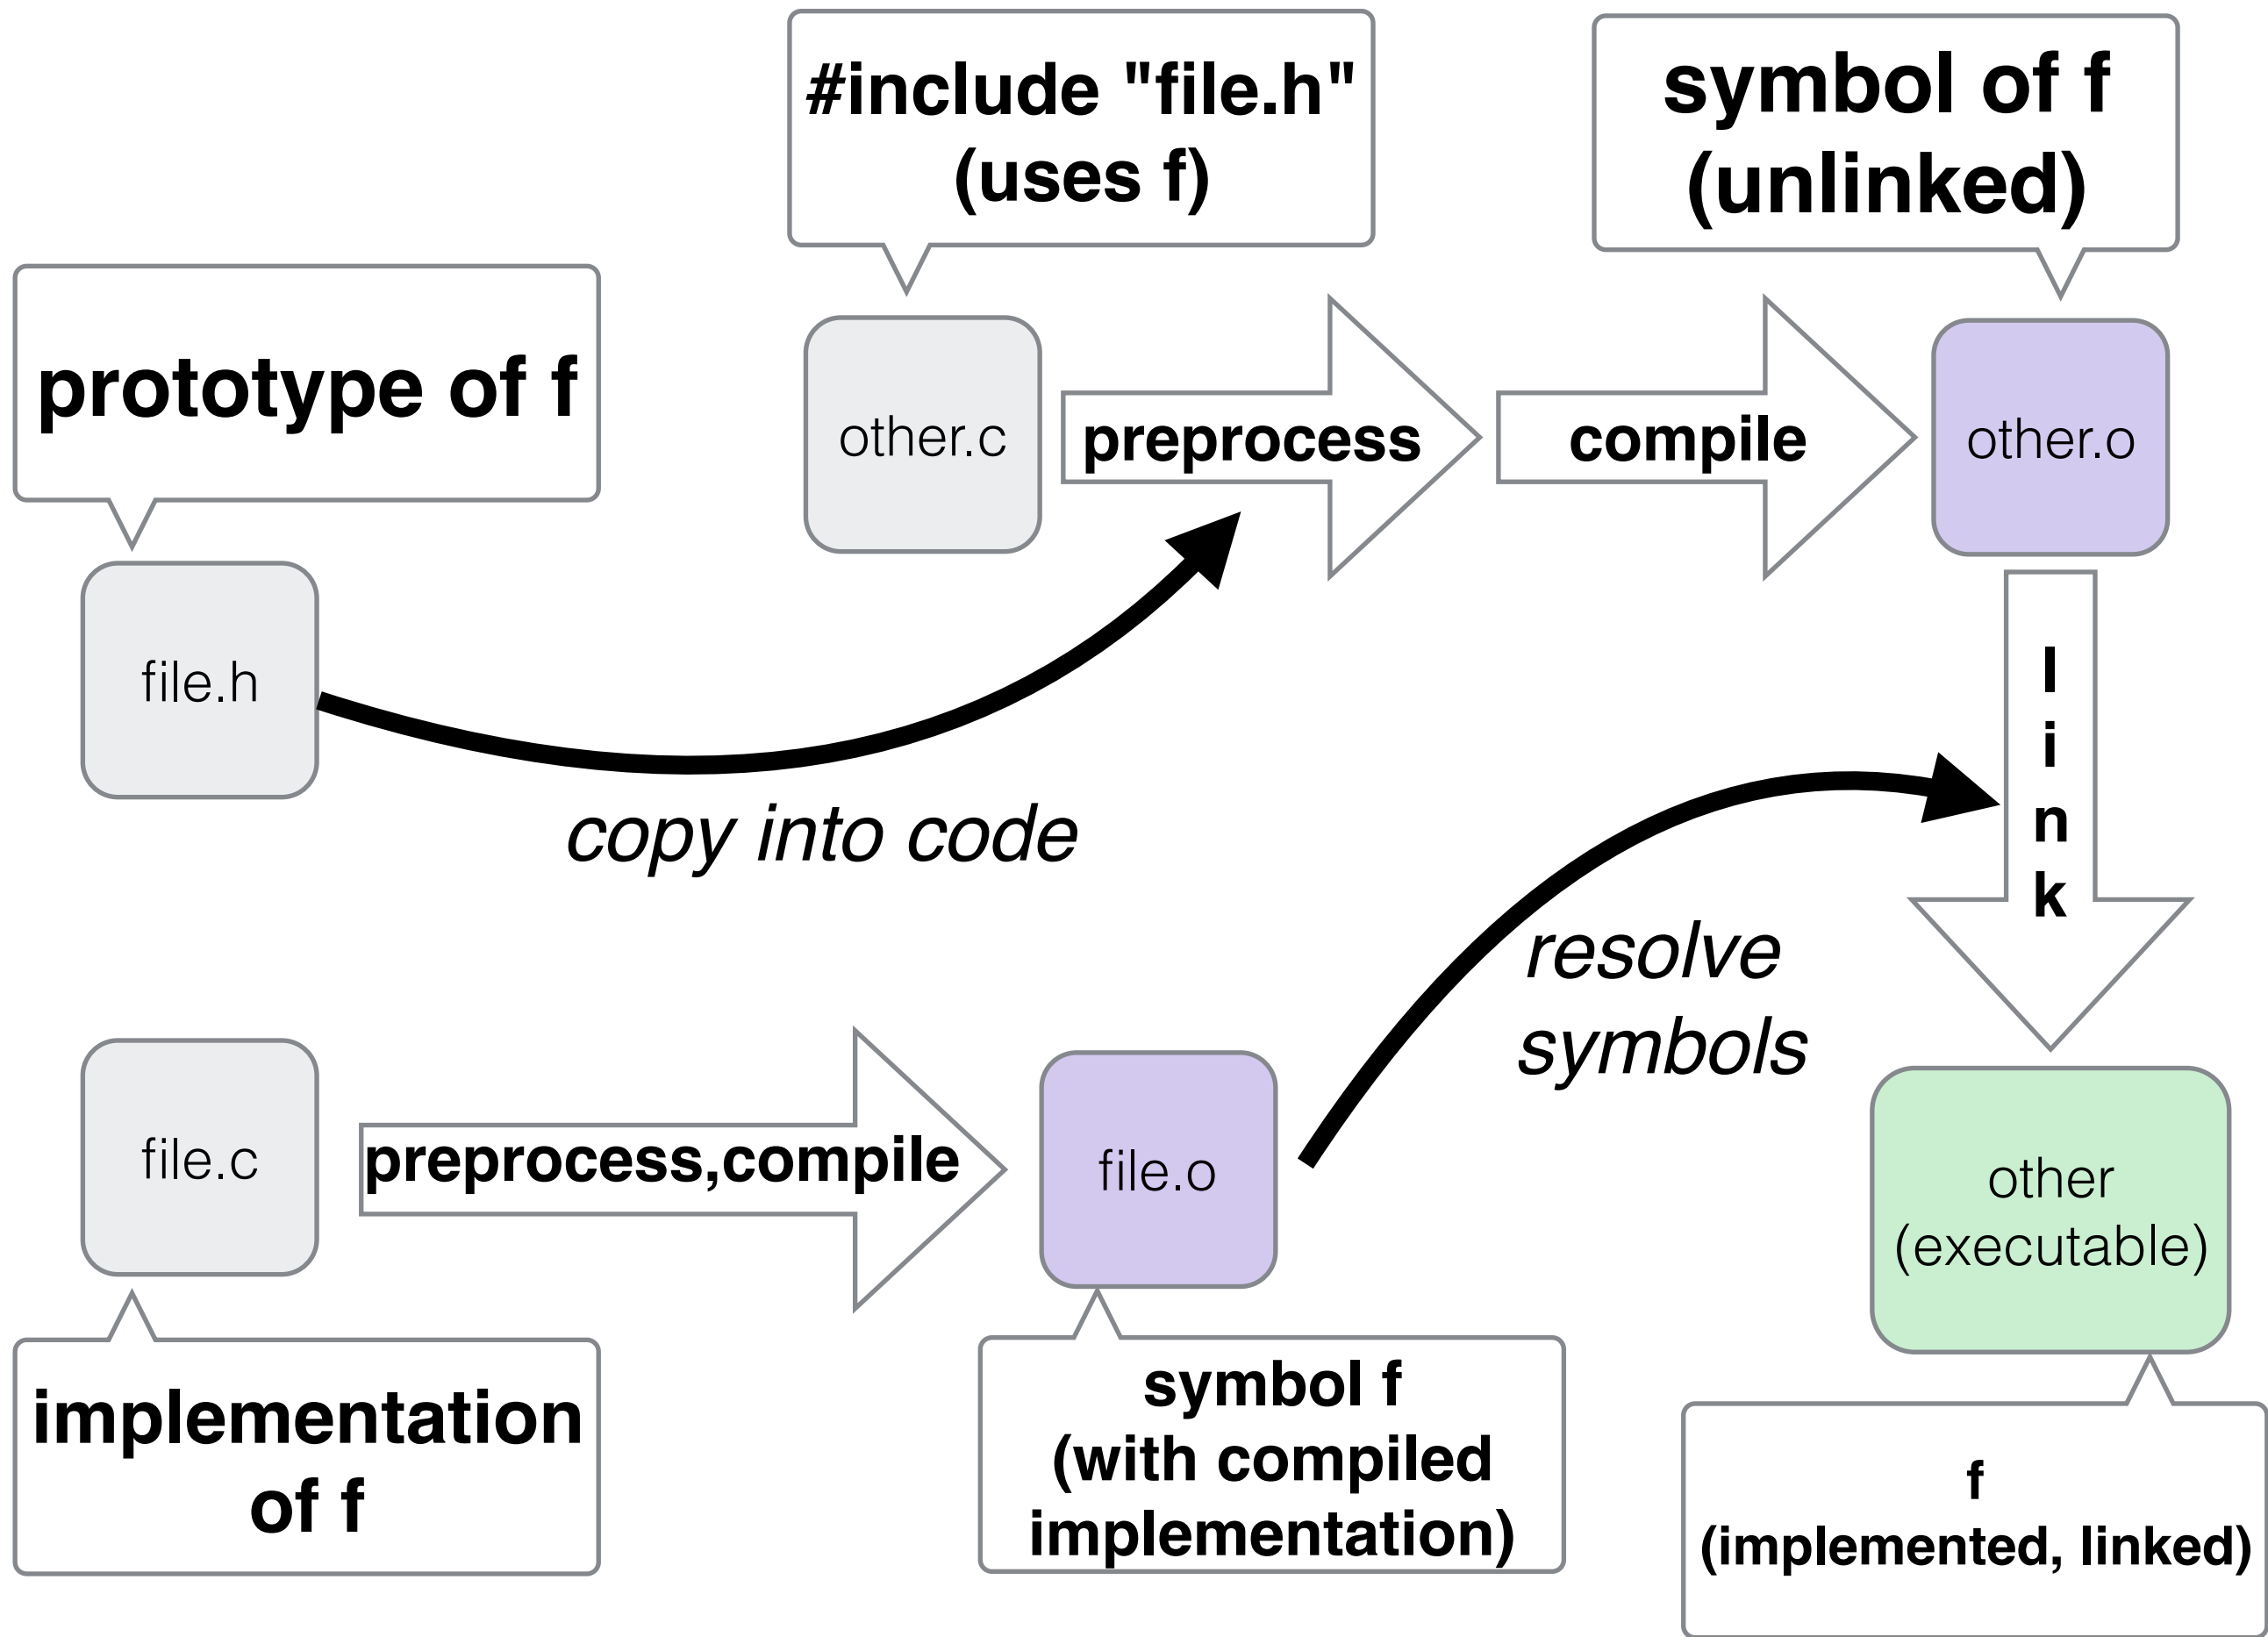
- Just as the preprocessor looks in "standard locations" for header files quoted in `< >`, the linker looks in standard locations for libraries.
- To add another location to the list to be searched, use `-L/path/to/new/location`

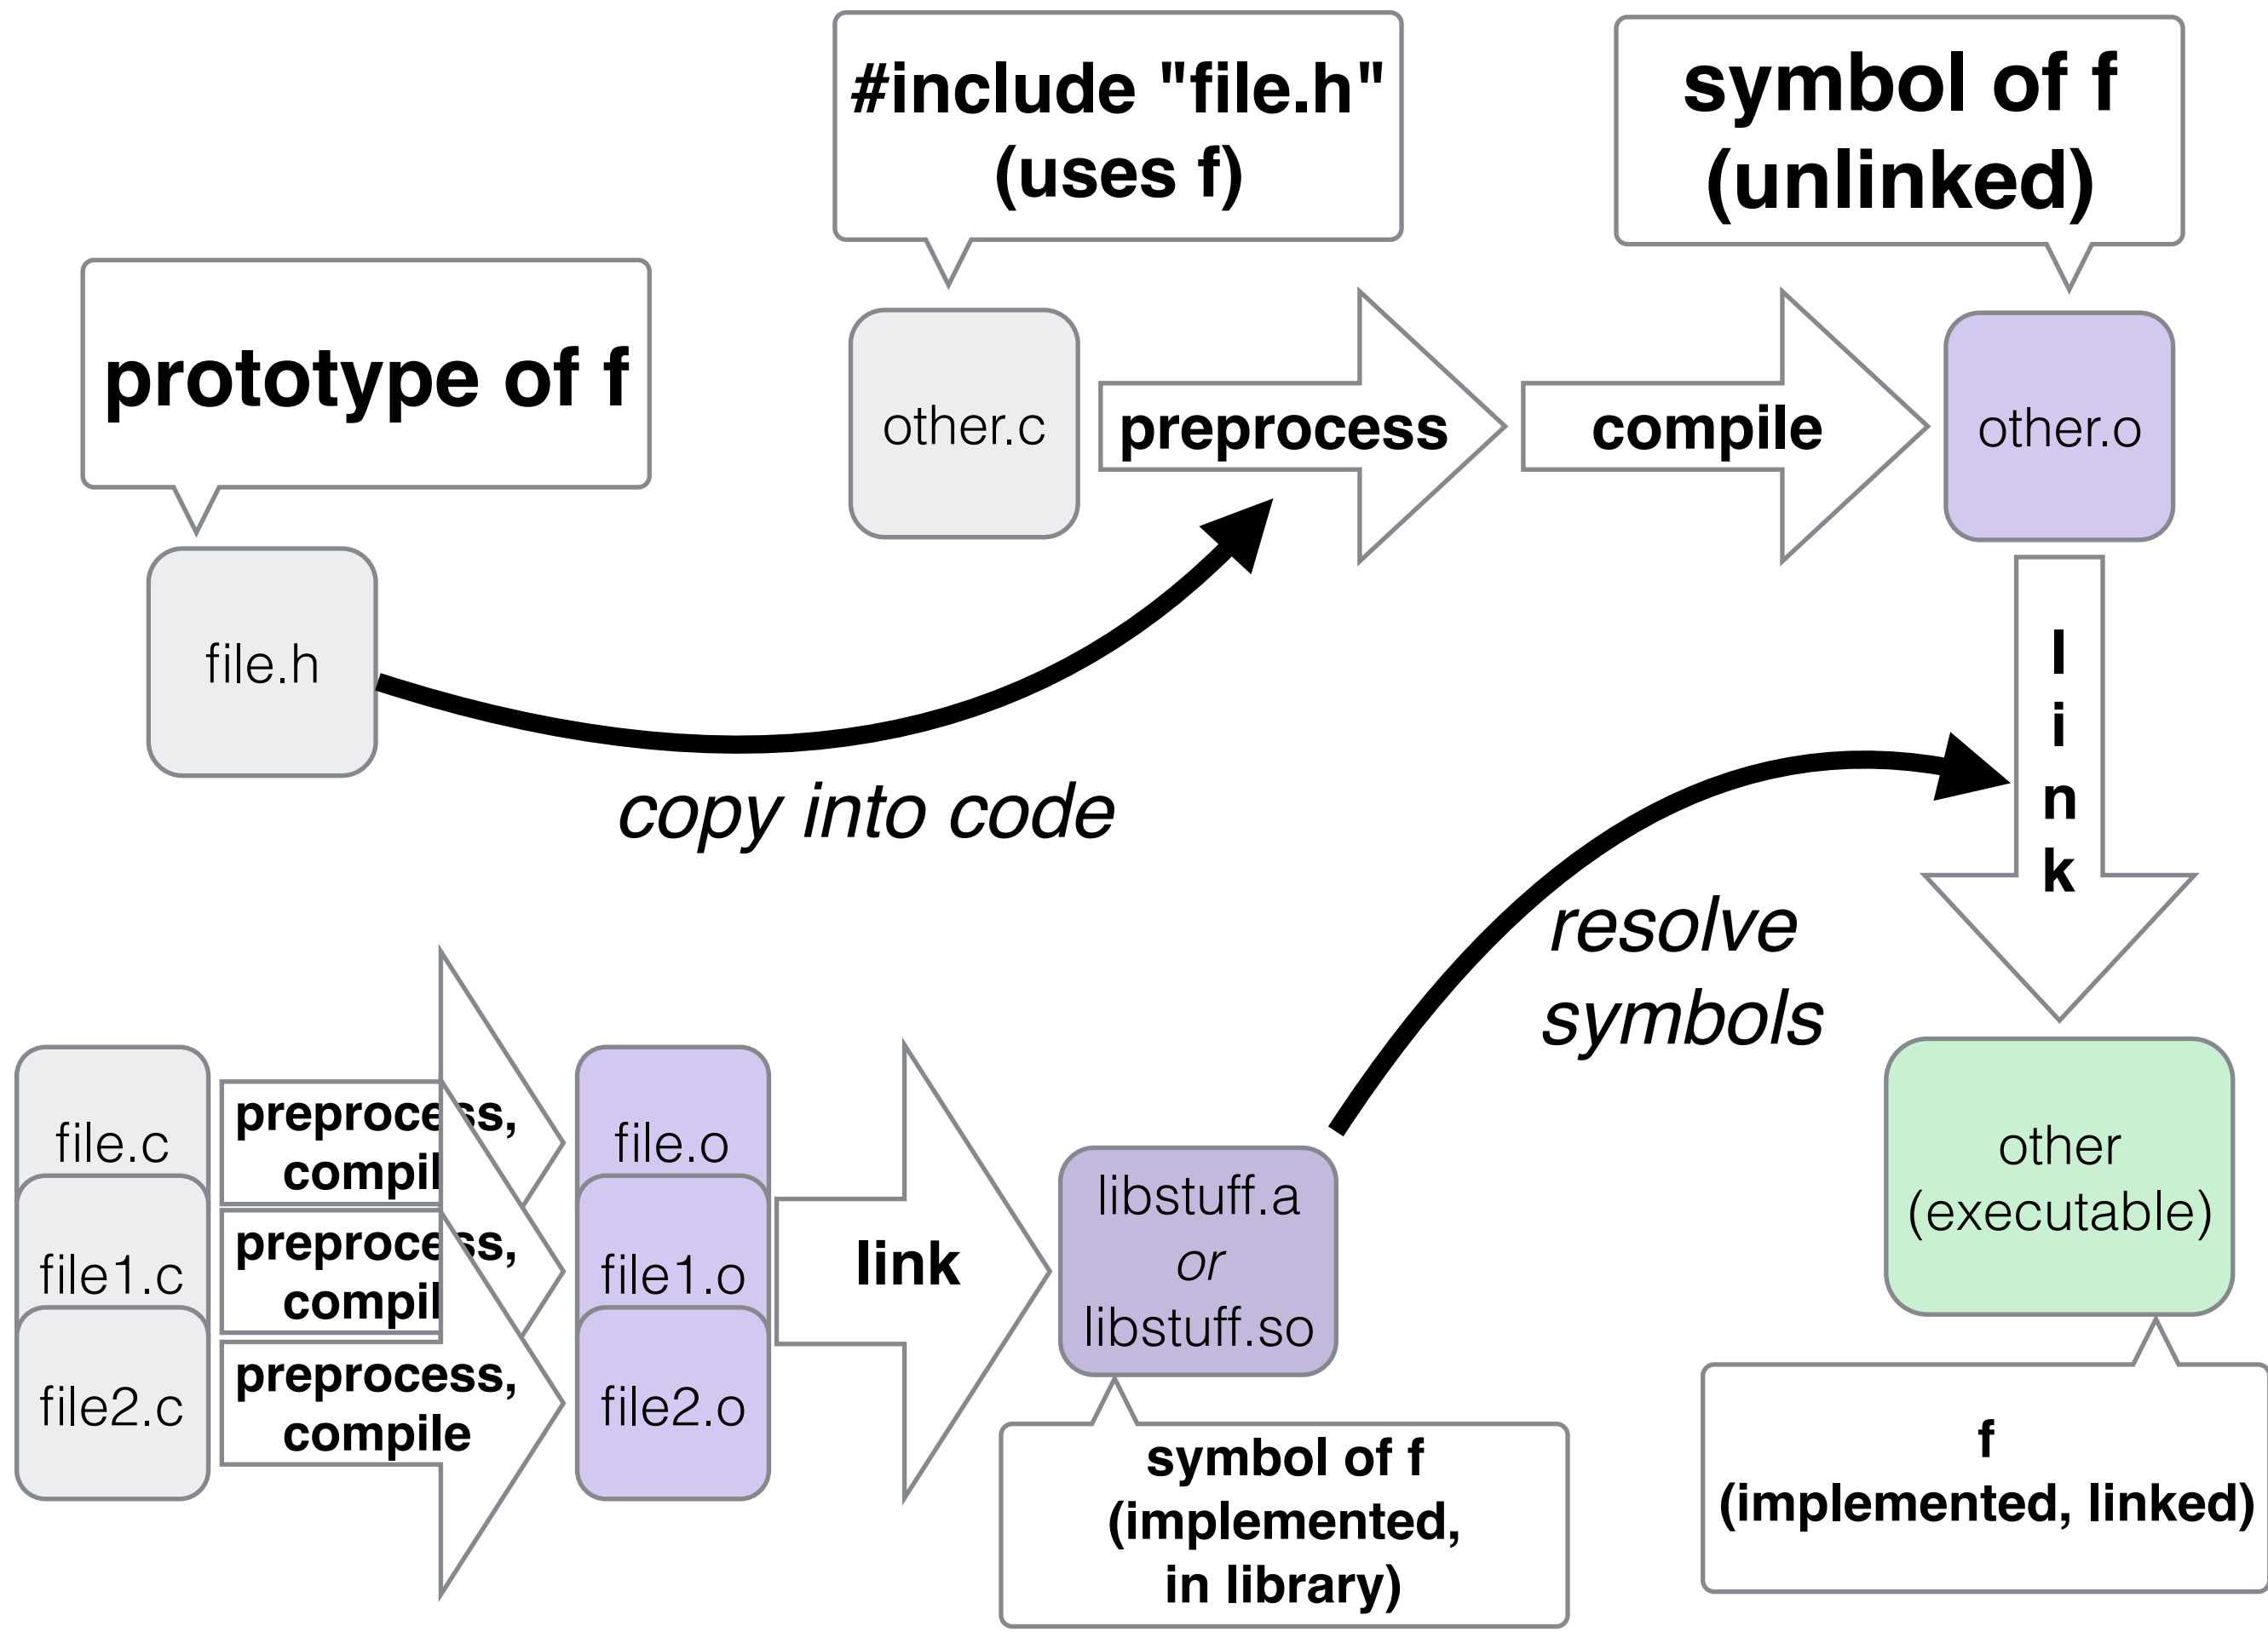
For example, to link a library with path

/home/myaccount/lib/libmylibrary.a

use

`-L/home/myaccount/lib/ -lmylibrary`





Libraries and Linking example

example.c

```
#include <stdio.h>
#include "function.h"

int main(void){
    printf("The result of function f is: %d\n", f(27));
    return 0;
}
```

function.h

```
//declaration of f (multiplies number by 2)
int f(int);
```

function.c

```
int f(int a) {
    return a*=2;
}
```

gcc example.c function.c
or (if .o files exist already)

gcc example.o function.o

nm a.out

(where does printf symbol
get resolved?)

ldd a.out

Libraries and Linking example

example.c

```
#include <stdio.h>
#include "function.h"

int main(void){
    printf("The result of function f is: %d\n", f(27));
    return 0;
}
```

replace with `#include "function2.h"`

Need to explicitly link libm.a/libm.so
(name is always the bit between "lib" and .a/.so)

function2.h

```
//function returns cosine of argument
double f(double);
```

```
gcc example.c function2.c -lm
```

```
ldd a.out
```

function2.c

```
#include <math.h>

double f(double a){
    //cos function declared in math.h, implemented in libm.a
    return cos(a);
}
```

The Mandelbrot set is
generated by:

for each point in the complex
plane, c


start with $z_0 = 0$

repeat:

$$z_{n+1} = z_n^2 + c$$

if z_n tends to infinity, colour
grey based on how fast it
does so.

if it doesn't, colour black (is in
the set)



This example is not examinable.

mandel.c

```
#include <stdio.h>
#include <complex.h>
#include <math.h>
#include <png.h>
#include <stdint.h>
#include "initpng.h"
```

complex.h header for complex numbers
math.h for complex math functions
png.h for writing png image
stdint.h for "exactly 8 bit wide" ints
initpng.h for own code to initialise a png

```
#define DIM 500
```

```
uint8_t mandel_kernel(double complex c);
```

```
int main(void) {
```

```
    uint8_t array[DIM][DIM];
    double complex c;
```

exactly 8 bit wide integer

floating point complex number!

```
    for(int i=0; i<DIM; i++){
        for (int j=0; j<DIM; j++){
            //align on range -2..2 in x and y
            c = (4*(double)i/DIM)-2 + ((4*(double)j/DIM)-2)*I;
            array[i][j] = mandel_kernel(c);
        }
    }
```

```
    FILE *fp = fopen("mandel.png", "wb");
```

```
    png_structp png = initpng(fp, DIM);
```

get "rows" of array, as ptrs

```
    uint8_t * rows[DIM];
    for(int i=0; i<DIM; i++) {
        rows[i] = (uint8_t *)array[i];
    }
```

```
    png_write_image(png, rows);
    png_write_end(png, NULL);
```

```
    fclose(fp);
```

```
}
```

mandelbrot set calculation

```
uint8_t mandel_kernel(double complex c) {
    uint8_t i = 100;
    double complex z = 0 + 0*I;
    while(cabs(z)<2.0 && i>0){
        z = cpow(z, 2) + c;
        i++;
    }
    return i;
}
```

***cpow raises z to the power 2
(efficiently)***
cabs gives the absolute value of z

gcc mandel.c initpng.o -lm -lpng

link to libm.so (for math)
link to libpng.so (for png creation)

Further reading

- You can read more about the preprocessor (and more features which are not covered in the course), here:
- <https://gcc.gnu.org/onlinedocs/cpp/>
- You can read a bit more about compiler optimisation here:
- <http://moodle2.gla.ac.uk/mod/page/view.php?id=149957>

References

- The documentation for the (GNU) implementation of the C Standard Library:
- <https://www.gnu.org/software/libc/manual/>
- (We have not covered a lot of the contents!)