

# C Lecture 7

File I/O

## File I/O

- To date, we've used various I/O functions without really exploring the whole subject of input-output.
- We're going to take a slightly more in-depth look at how C does input and output here.
- Our context is "writing and reading to files"...
  - but *everything* we do I/O to in Linux is a file (including keyboard and terminal), so this also applies to the other I/O we've done.
- All I/O functions are declared in `stdio.h` (and implemented in `libc`)

## "Stream I/O"

- In this course, we will consider two kinds, of "stream I/O" ("text" and "binary").
- You can think of a stream as a logical interface between your program and a "file".
- Functions can stream characters out of a file (reading) or stream them into a file (writing).
- A stream can connect your program to any file (or file-like thing - such as the special `STDIO`, `STDIN`, `STDERR` files for handling user input/output).

## Text Stream I/O

This is some text that is in the input stream.\n More lines of text...

fgets(...)

fgets takes characters from the stream, until it reaches a newline

This is some text that is in the input stream.\n More lines of text...



After fgets returns, the stream "no longer contains" the characters it took... If the stream represents a file, then the file itself isn't changed – but we've "moved on in the file" past those characters.

# ASCII

- The "characters" in a text stream are encoded in the format known as ASCII (American Standard Code for Information Interchange).
- This is also what is used for your `char` variables.
- It assigns values from 0 to 127 to different symbols (including the Latin letters, lowercase and uppercase).
- A Text I/O stream, therefore, is a sequence of ASCII symbols.

*ASCII cannot represent non-Latin alphabets. More modern languages than C use other codes – Unicode, for example – which can represent many more characters, alphabets, etc.*

ASCII: American Standard Code for Information Interchange.

NULL = \0 character				space character											
Dec	Hx	Oct	Chr	Dec	Hx	Oct	Chr	Dec	Hx	Oct	Chr	Dec	Hx	Oct	Chr
0	0	000	NULL (null)	32	20	040	Space	64	40	100	@#64: B	96	60	140	@#96: a
1	1	001	SOH (start of heading)	33	21	041	!	65	41	101	@#65: A	97	61	141	@#97: b
2	2	002	STX (start of text)	34	22	042	"	66	42	102	@#66: B	98	62	142	@#98: c
3	3	003	ETX (end of text)	35	23	043	#	67	43	103	@#67: C	99	63	143	@#99: d
4	4	004	EOF (end of transmission)	36	24	044	\$	68	44	104	@#68: D	100	64	144	@#100: e
5	5	005	ENQ (enquiry)	37	25	045	%	69	45	105	@#69: E	101	65	145	@#101: f
6	6	006	ACK (acknowledge)	38	26	046	&	70	46	106	@#70: F	102	66	146	@#102: g
7	7	007	BEL (bell)	39	27	047	'	71	47	107	@#71: G	103	67	147	@#103: h
8	8	010	BS (backspace)	40	28	050	(	72	48	110	@#72: H	104	68	150	@#104: i
9	9	011	TAB (horizontal tab)	41	29	051	)	73	49	111	@#73: I	105	69	151	@#105: j
10	A	012	LF (NL line feed, new line)	42	2A	052	*	74	4A	112	@#74: J	106	70	152	@#106: k
11	B	013	VT (vertical tab)	43	2B	053	+	75	4B	113	@#75: K	107	71	153	@#107: l
12	C	014	FF (NP form feed, new page)	44	2C	054	,	76	4C	114	@#76: L	108	72	154	@#108: m
13	D	015	CR (carriage return)	45	2D	055	-	77	4D	115	@#77: M	109	73	155	@#109: n
14	E	016	SO (shift out)	46	2E	056	.	78	4E	116	@#78: N	110	74	156	@#110: o
15	F	017	SI (shift in)	47	2F	057	/	79	4F	117	@#79: O	111	75	157	@#111: p
16	10	020	DLE (data link escape)	48	30	060	0	80	50	120	@#80: P	112	76	160	@#112: q
17	11	021	DC1 (device control 1)	49	31	061	1	81	51	121	@#81: Q	113	77	161	@#113: r
18	12	022	DC2 (device control 2)	50	32	062	2	82	52	122	@#82: R	114	78	162	@#114: s
19	13	023	DC3 (device control 3)	51	33	063	3	83	53	123	@#83: S	115	79	163	@#115: t
20	14	024	DC4 (device control 4)	52	34	064	4	84	54	124	@#84: T	116	80	164	@#116: u
21	15	025	NAK (negative acknowledge)	53	35	065	5	85	55	125	@#85: U	117	81	165	@#117: v
22	16	026	SYN (synchronous idle)	54	36	066	6	86	56	126	@#86: V	118	82	166	@#118: w
23	17	027	ETB (end of trans. block)	55	37	067	7	87	57	127	@#87: W	119	83	167	@#119: x
24	18	030	CAN (cancel)	56	38	070	8	88	58	130	@#88: X	120	84	170	@#120: y
25	19	031	EM (end of medium)	57	39	071	9	89	59	131	@#89: Y	121	85	171	@#121: z
26	1A	032	SUB (substitute)	58	3A	072	:	90	5A	132	@#90: Z	122	86	172	@#122: {
27	1B	033	ESC (escape)	59	3B	073	;	91	5B	133	@#91: [	123	87	173	@#123:
28	1C	034	FS (file separator)	60	3C	074	<	92	5C	134	@#92: \	124	78	174	@#124: }
29	1D	035	GS (group separator)	61	3D	075	=	93	5D	135	@#93: ]	125	79	175	@#125: ~
30	1E	036	RS (record separator)	62	3E	076	>	94	5E	136	@#94: ^	126	76	176	@#126: _
31	1F	037	US (unit separator)	63	3F	077	?	95	5F	137	@#95: `	127	77	177	@#127: DEL

## Opening a File

```
#include <stdio.h>

int main(void)
{
    char mytext[100];

    FILE *file_ptr;
    file_ptr = fopen("myFile.txt", "r");
    fgets(mytext, sizeof(mytext), file_ptr);
    puts(mytext);
    return 0;
}
```

FILE is a special structured type which represents the stream of data flowing to or from a file.

fopen(char[] filename, char[] mode) returns a pointer to type file

The mode specifies how we are going to access the file (read, write, append...)

The special file pointers STDOUT, STDIN and STDERR are automatically opened when a program starts and correspond to the same concepts you met in the Bash part of the course.

## File access modes

File mode	Intent	Equivalent Bash operator	Effect of adding a +
"r"	Just read the file	<	"r+" - read and write the file (from the start)
"w"	Just overwrite the file	>	"w+" - overwrite the file, but allow reading too
"a"	Just write stuff to end of file	>>	"a+" - start writing to end of file (but allow reading too)

Truncates (erases contents!) if it opens an existing file. Otherwise creates a new file

```
FILE * fp = fopen("myfile.txt", "r");
```

## Reading from a File - fgets

- char \* fgets(char \* string, int len, FILE \* file);

```
#include <stdio.h>

int main(void)
{
    char mytext[100];
    char *p1;

    FILE *file_ptr;
    file_ptr = fopen("myFile.txt", "r");
    p1 = fgets(mytext, sizeof(mytext), file_ptr);
    puts(mytext);
    puts(p1);
    // puts(fgets(mytext, sizeof(mytext), file_ptr));
    return 0;
}
```



## Reading from a File - fscanf

- int fscanf(FILE \* file, char \* format, ...)

Any number of pointers to variables to match format

- fscanf returns an int, which is the number of variables it successfully put values into.
- (This can be less than the number requested, if it was unable to interpret some of the stream in the requested way.)
- fscanf returns EOF if it reaches the end of a file (and therefore there's nothing more to read).

fscanf(stdin, "%d is an integer\n", &myint) is identical to scanf("%d is an integer\n", &myint)

## General I/O Functions

- In general, I/O functions are written in the form **object**verb**subject**.  
 Where **object** can be **f**, for files, **s**, for strings, or omitted, for "default/terminal".  
 (There are other options, such as **sn** for "string, along with a max length", but we leave those for the documentation.)
- And **subject** is **f** for "formatted values", **s** for "to string", or **c** for "single character".
- `scanf()`, `sscanf()`, `printf()`, `fprintf()`, `getc()`
- `scanf()`, e.g., uses a **format string** to convert values to text, and stores the result in a **string**.
- (There are some violations of these general rules, especially for the get and put families, and some combinations do not exist, so check the documentation before using.)

## fscanf() Example

```
#include <stdio.h>

int main(void)
{
    char mytext[100];
    int i = 0, j = 0;
    float a = 0.0;

    FILE *file_ptr;
    file_ptr = fopen("myFile.txt", "r");
    puts(fgets(mytext, sizeof(mytext), file_ptr));
    puts(fgets(mytext, sizeof(mytext), file_ptr));
    fscanf(file_ptr, "%d %f\n", &i, &a);
    // j = fscanf(file_ptr, "%d %f\n", &i, &a);

    printf("Read in %d %f\n", i, a, j);

    return 0;
}
```

## fscanf

900 6 8.000\n980 h 8.100\n577 86 6.456\n8 77 8.100\n900 6 8.000\n

900 6 8.000\n980 h 8.100\n577 86 6.456\n8 77 8.100\n900 6 8.000\n

`fscanf(fp, "%d %d %f\n", ...)` format matches next characters in stream  
(`fscanf` returns 3)

980 h 8.100\n577 86 6.456\n8 77 8.100\n900 6 8.000\n7 34 4.4359\n

`fscanf(fp, "%d %d %f\n", ...)` format can't match second conversion specifier  
(`'h'` is not interpretable as an int)

h 8.100\n577 86 6.456\n8 77 8.100\n900 6 8.000\n7 34 4.4359\n789

stream is left at point of first "non-matching" characters.  
`fscanf` returns 1;  
 first variable (pointer) passed to `fscanf` is assigned value 980 (as an int),  
 the other two are *unchanged*.

## Writing to a File

- `int fputs(char * string, FILE * file);`
- `int fprintf(FILE * file, char * format, ...);`  
 Any number of values to match format
- `fputs` and `fprintf` both return an int - for `fprintf`, this is the number of characters it wrote.
- If they fail to write to the file, they instead return the special value EOF.
- Unlike `puts`, `fputs` does not add a '\n' to output.

```
fprintf(file, "%d is an integer\n", myint);
```

`fprintf(stdout, "%d is an integer\n", myint)` is identical to `printf("%d is an integer\n", myint)`

## Streams and Buffers

- Strictly, writes to a stream are not immediately reflected in the file they are connected to.
- (Physical media like hard disks, or optical disks, takes real time to write stuff.)
- Instead, the writes are "queued up" in a buffer of things needing to be written.
- Periodically, the buffer is emptied, and its contents actually committed to the file.

## Closing a File

- `int fflush(FILE * fp);`
- `fflush` makes sure that all the data we've written so far has actually been committed to the file itself (it flushes the buffer immediately).
- `int fclose(FILE * fp);`
- `fclose` does a flush, and then removes the connection between `fp` and the file it's attached to. If successful, it returns 0.
- After closing a file, the file pointer *cannot* be used for I/O until assigned a new file with an `open`.

**NEVER** `fclose stdin, stdout or stderr!` (`fflush` can be useful to ensure stuff is printed immediately).

## Text Stream I/O Example

```
#include <stdio.h>

int main(void)
{
    FILE *fp;
    fp = fopen("testfile.txt", "w");

    fprintf(fp, "Number %d %f", 5, 5.0 );

    fflush(fp); //ensure everything written - not needed in this example
    fclose(fp); //fclose flushes buffer before closing file

    fp = fopen("testfile.txt", "r");

    int a;
    float b;

    fscanf(fp, "Number %d %f", &a, &b);
    fclose(fp);

    printf("Float in file was %f\n", b );

    return 0;
}
```

mac-db2016:lecture7 Dave\$ hexdump -C testfile.txt  
 00000000 |4e 75 64 62 65 72| 20 35 20 35 2e 30 30 30 30 30 |Number 5 5.000001  
 00000010 |30 | 00  
 00000011 |

## Binary Stream I/O

- Rather than opening a file and streaming the contents as if it were text, we may want to stream the bytes to/from a file directly.
- A file created this way is called a binary file. Bits in text files represent ASCII characters, but the bits in binary files represent the actual custom data.
- Binary stream I/O can be more efficient than text stream I/O
- Also doesn't lose precision of floating point values.
- But you can't just read a binary file because you don't know how to interpret the contents (is the first byte part of an integer? Or a character? Or what?)

I/O type	Text Stream	Binary Stream
Pros	Human readable	Exact representation of variables
Cons	Precision loss Often larger than needed	Opaque Less portable

## Opening, Closing Binary Streams

- We can still use `fopen` and `fclose` to work with Binary Stream I/O.
- There's just one change: the file mode must have a "b" added to it, to ensure we directly read precisely what's in the file.
- So, "r" -> "rb", "w+" -> "wb+" and so on.

```
FILE * fp = fopen("myfile.txt", "rb");
```

## Reading, Writing Binary

- `long fread(void * start, long size, long num, FILE * file);`
- `long fwrite(void * start, long size, long num, FILE * file);`
- `fread` reads **size\*num** bytes from **file**, and inserts them into memory starting at **start**. (It returns how many it actually wrote).
- `fwrite` does the same thing, but takes *from* **start**, and writes *into* **file**.

```
int a[5];
fread(a, sizeof(int), 5, fp);
```

## Binary Stream I/O Example

```
#include <stdio.h>
#include <string.h>

int main(void)
{
    FILE *fp;
    fp = fopen("testfile.bin", "wb");

    char str[] = "Example String";
    int a[5] = {0,1,2,3,4};

    fwrite(str, sizeof(char), strlen(str), fp);
    fwrite(a, sizeof(int), 5, fp);
    fclose(fp);

    char str2[] = "00000000000000";
    int b[5] = {0,0,0,0,0};

    fp = fopen("testfile.bin", "rb");
    fread(str2, sizeof(char), strlen(str2), fp);
    fread(b, sizeof(int), 5, fp);

    printf("Read in %s %d,%d,%d,%d,%d\n", str2, b[0], b[1], b[2], b[3], b[4]);

    return 0;
}
```