



Why Linux?

Dr. Gareth Roy (x6439)
gareth.roy@glasgow.ac.uk

- A brief history lesson
- Why Linux?
- The parts of an Operating System
- The Command Line

- A brief history lesson
- Why Linux?
- The parts of an Operating System
- The Command Line

IBM

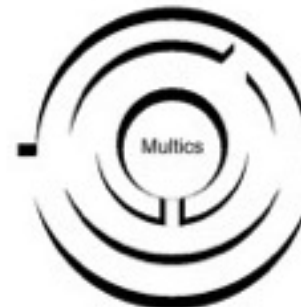
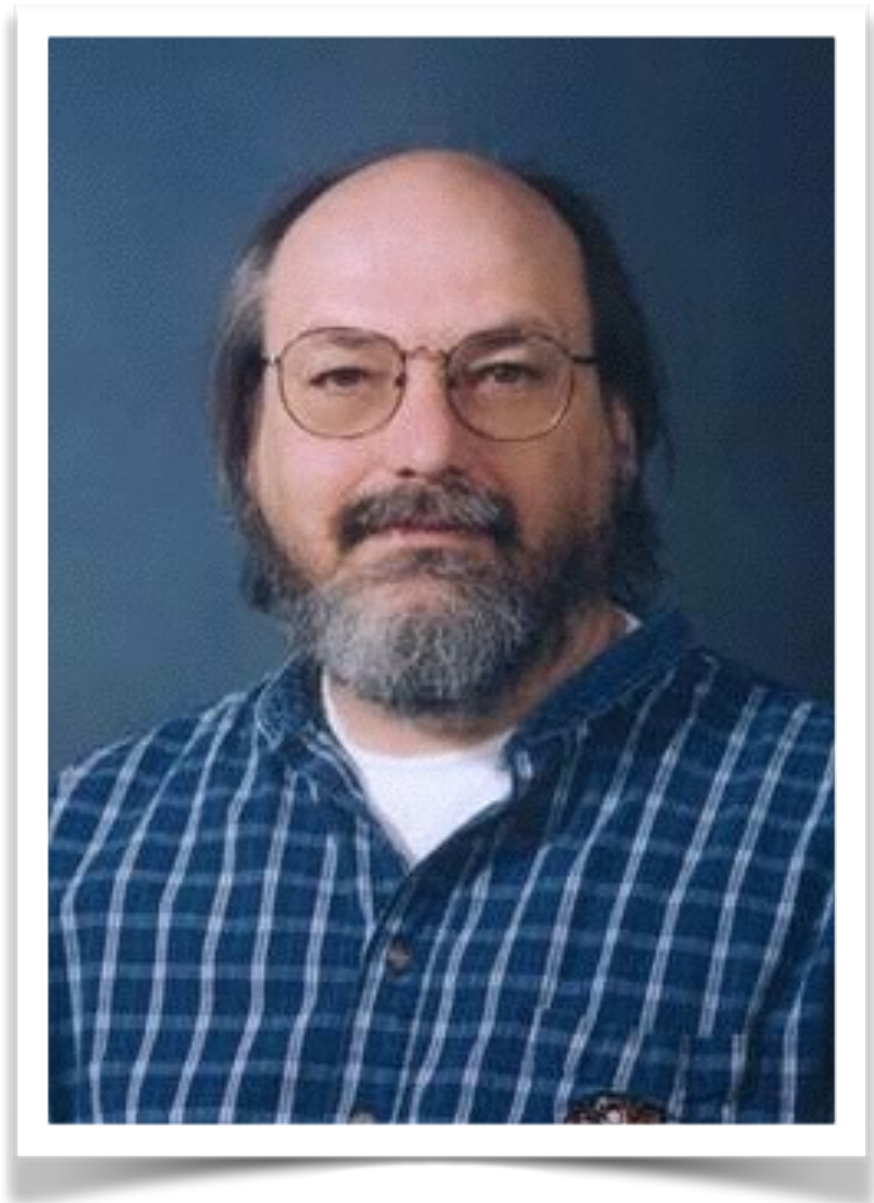
WARD SHERATON

IBM DATACENTRE

THE B&N

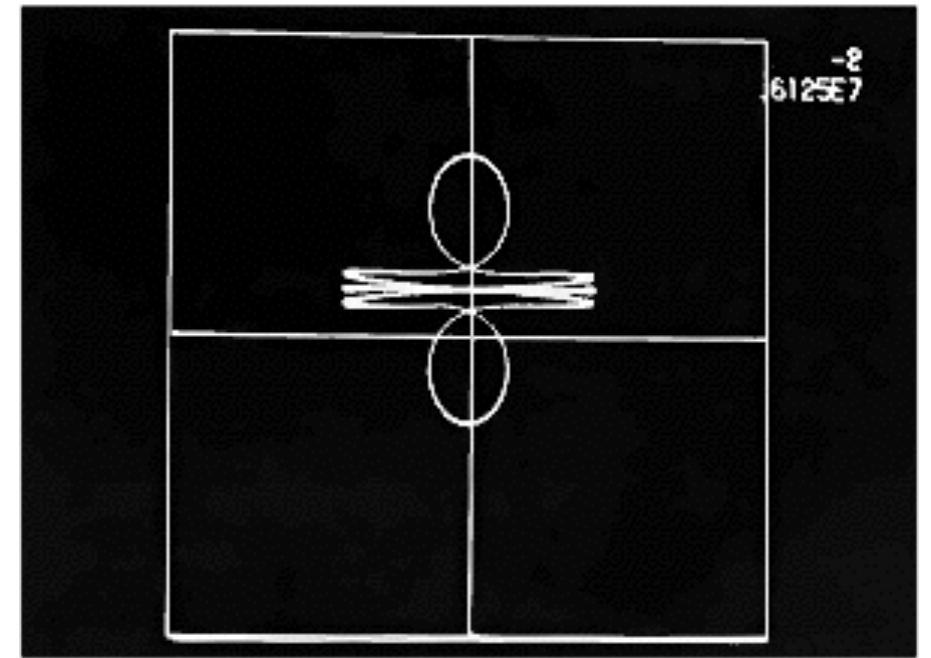
Kenneth Lane Thomson (1943 -)

- Received his Bachelor's and Master's degrees in E&EE from the University of California at Berkley
- In 1966 was hired by Bell Labs (owned by AT&T) to work on the MULTICS project.
- MULTICS (**M**ultiplexed **I**nformation and **C**omputing **S**ervice) was a mainframe time sharing operating system.
- Worked on developing the OS until Bell Labs withdrew from the consortium in 1969.



Space Travel

- While working on MULTICS, Ken created a game called Space Travel.
- It allowed the player to travel around a 2D solar system.
- Originally intend for MULTICS, Ken re-wrote it to run on an obsolete PDP-7 located in the lab.
- Ken wanted a better system to run his game on so in the Summer of 1969 he took one month to write:
 - Kernel
 - Shell
 - Editor
 - Assembler



Birth of UNIX

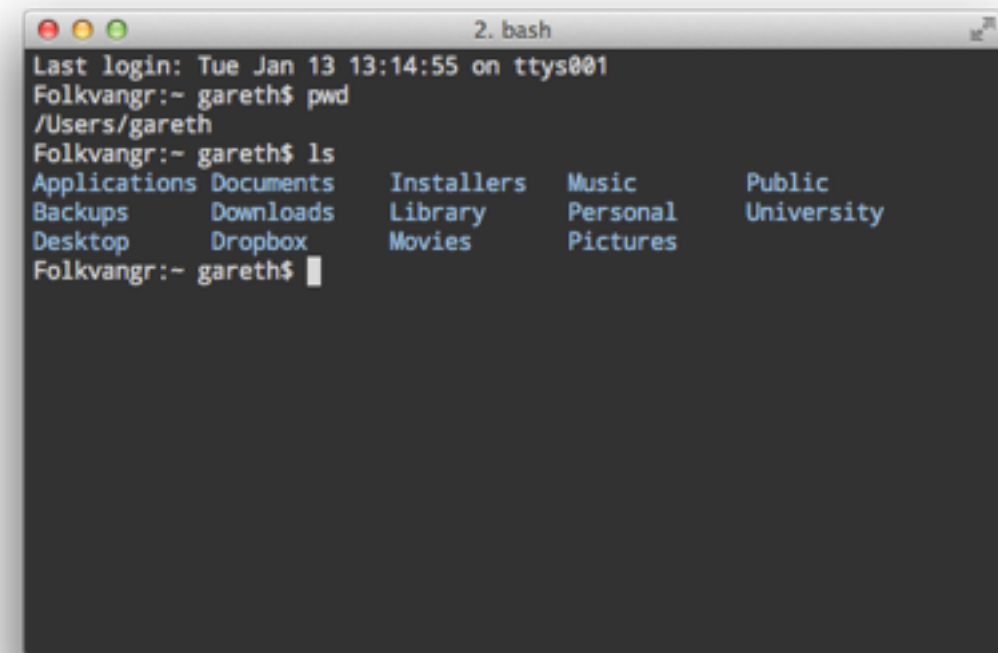
- Bell Labs left the MULTICS consortium in 1969.
- Bell Labs purchased new PDP-11, and development continued on what would become UNIX.
- In 1970 Ken created the 'B' programming language which was the precursor to C.
- Joined by Dennis Ritchie in 1972, Ken re-wrote the UNIX kernel in C.
- UNIX was presented to the world in 1973.



Ken Thomson and Dennis Ritchie working on the PDP-11 used to develop UNIX

What is UNIX

- UNIX is a Multi-user, Multi-tasking operating system (at present a family of Operating Systems).
- Any UNIX system has some key features:
 - the use of plain text for storing data
 - a hierarchical file system
 - devices and processes represented as files
 - the use of a large number of software tools, small programs that can be composed as opposed to using a single monolithic program
- In 2015 to be called a UNIX requires certification by the OpenGroup

A screenshot of a terminal window titled '2. bash'. The window shows a login session for 'gareth' on 'ttys001'. The user enters 'pwd' and the output is '/Users/gareth'. Then the user enters 'ls' and the output is a list of directories: Applications, Documents, Installers, Music, Public, Backups, Downloads, Library, Personal, University, Desktop, and Dropbox. The prompt 'Folkvangr:- gareth\$' is visible at the end of the last line.

<http://www.unix.org/version4/>

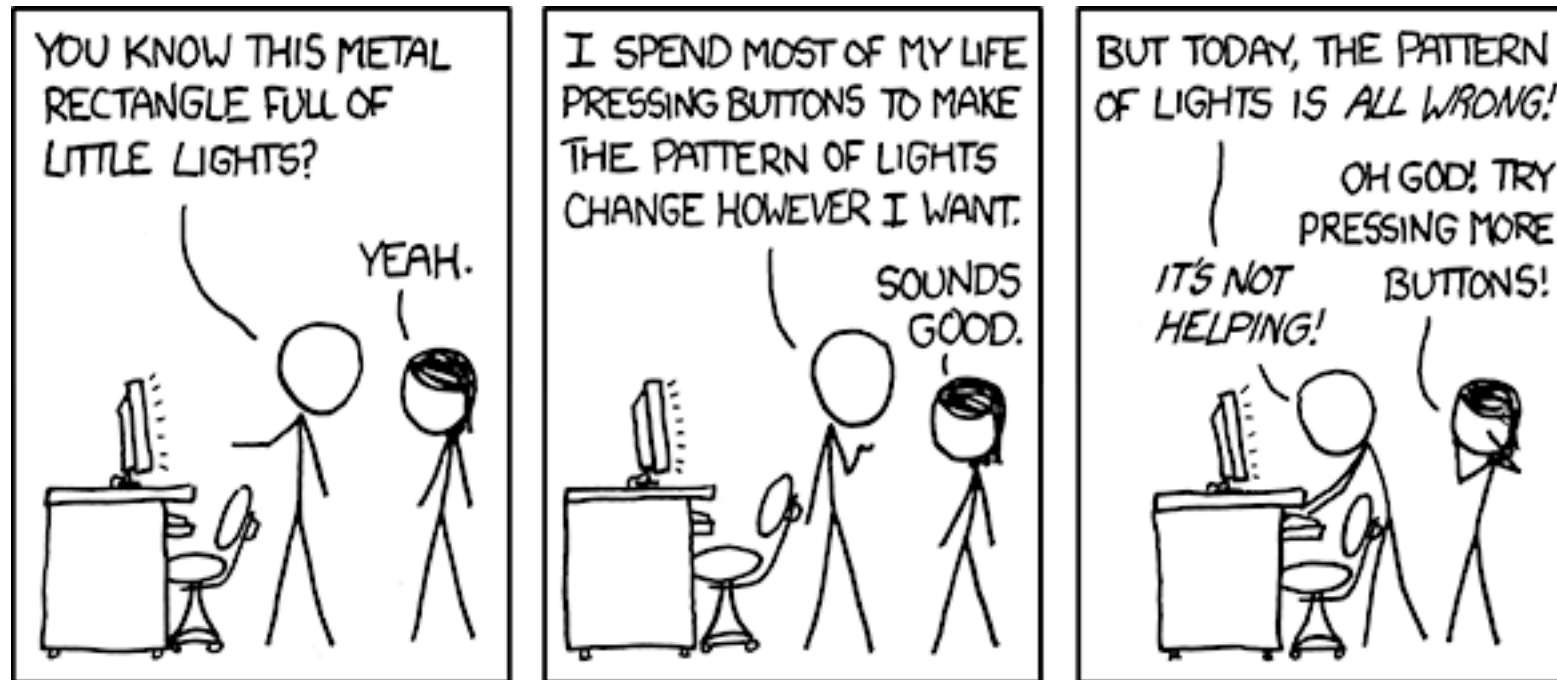
Linus Benedict Torvalds (1969 -)

- After becoming frustrated with the License associated with Minix, Linus began work on his own operating system.
- In 1991, Linux v.0.01 was released on comp.os.minix
- Originally intended to be called Freax, renamed by a co-worker (Ari Lemmke) when it was place on the server.
- Linux was only the OS Kernel itself and so a complete Linux OS heavily leveraged the GNU utilities and tools.
- This later led to Linux being released under a GPL license becoming open source.



"I'm doing a (free) operating system (just a hobby, won't be big and professional like gnu) for 386(486) AT clones."

- A brief history lesson
- Why Linux?
- The parts of an Operating System
- The Command Line



PARENTS: TALK TO YOUR KIDS ABOUT LINUX... BEFORE SOMEBODY ELSE DOES.

Why Linux?

- Linux is a free, with no cost required to obtain a complete fully featured OS.
- Linux source code is freely available and can be modified to fit user needs.
- Linux is a multi-user, multi-tasking operating system based on the tested principles of UNIX.
- Linux is ubiquitous in scientific computing with 485 of the top 500 supercomputers in the world running Linux.
- Scientific code bases for Cosmology, Particle Physics, Nuclear, Atmospheric, Fluid Dynamics all run on Linux platforms.

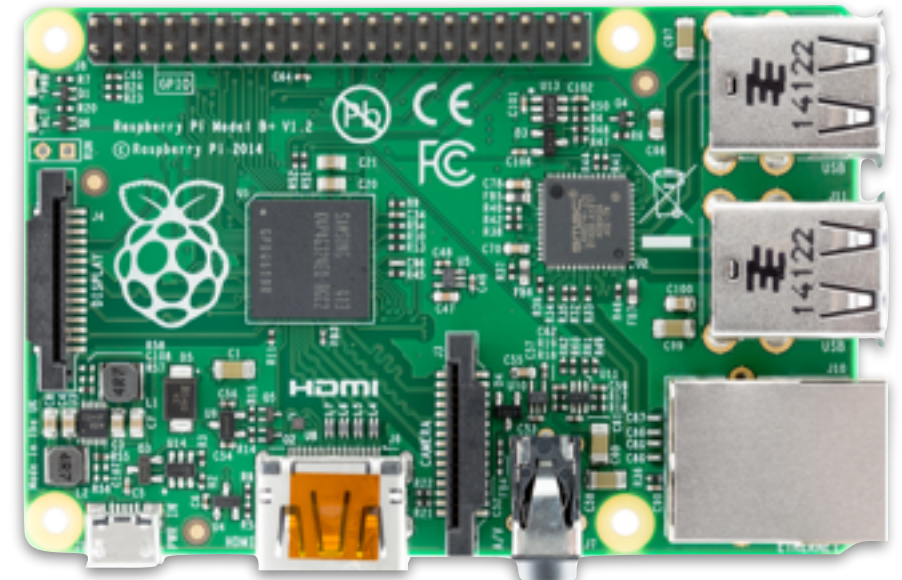


Tianhe-2 (MilkyWay-2)

- 3,120,000 CPU Cores
- 1,024,000 GB of RAM
- 33,862.7 TFlop/s
- 17,808 kW

Why Linux?

- Due to its open nature Linux has been ported to many platforms and architectures.
- It scales from embedded devices such as the Raspberry PI to powering companies such as Google and Facebook.
- Linux is on your phone (Android), it's on your router (dd-wrt), it's on your TV (lots), in your car (BMW), it runs the city of Munich, and your data is already stored there if you use Facebook/Google/DropBox etc...
- It runs much of the worlds internet services with, until recently, most deployments running LAMP.



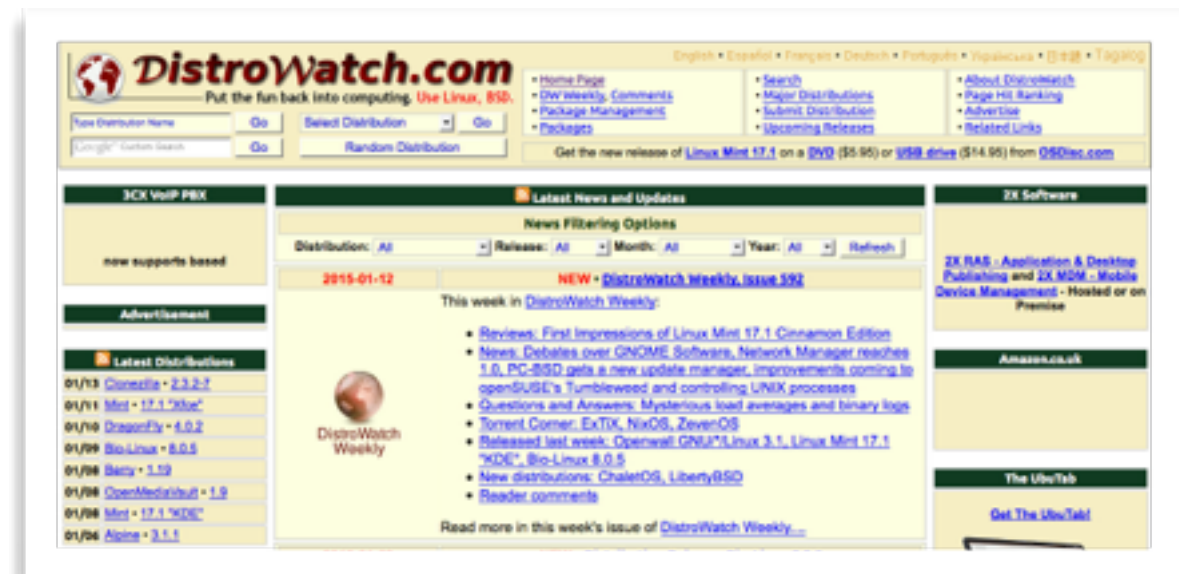
Why NOT Linux?

- Rough around the edges, oft times documentation is poor and no single source for information.
- Driver support can be weak for new hardware, and particularly for hardware that has closed data sheets (WiFi).
- Many applications target other platforms, so no Photoshop, MS Office or Assassin's Creed.
- Many problems need to be resolved on the command line rather than through more intuitive wizards.



So I've convinced you, how do you get it?

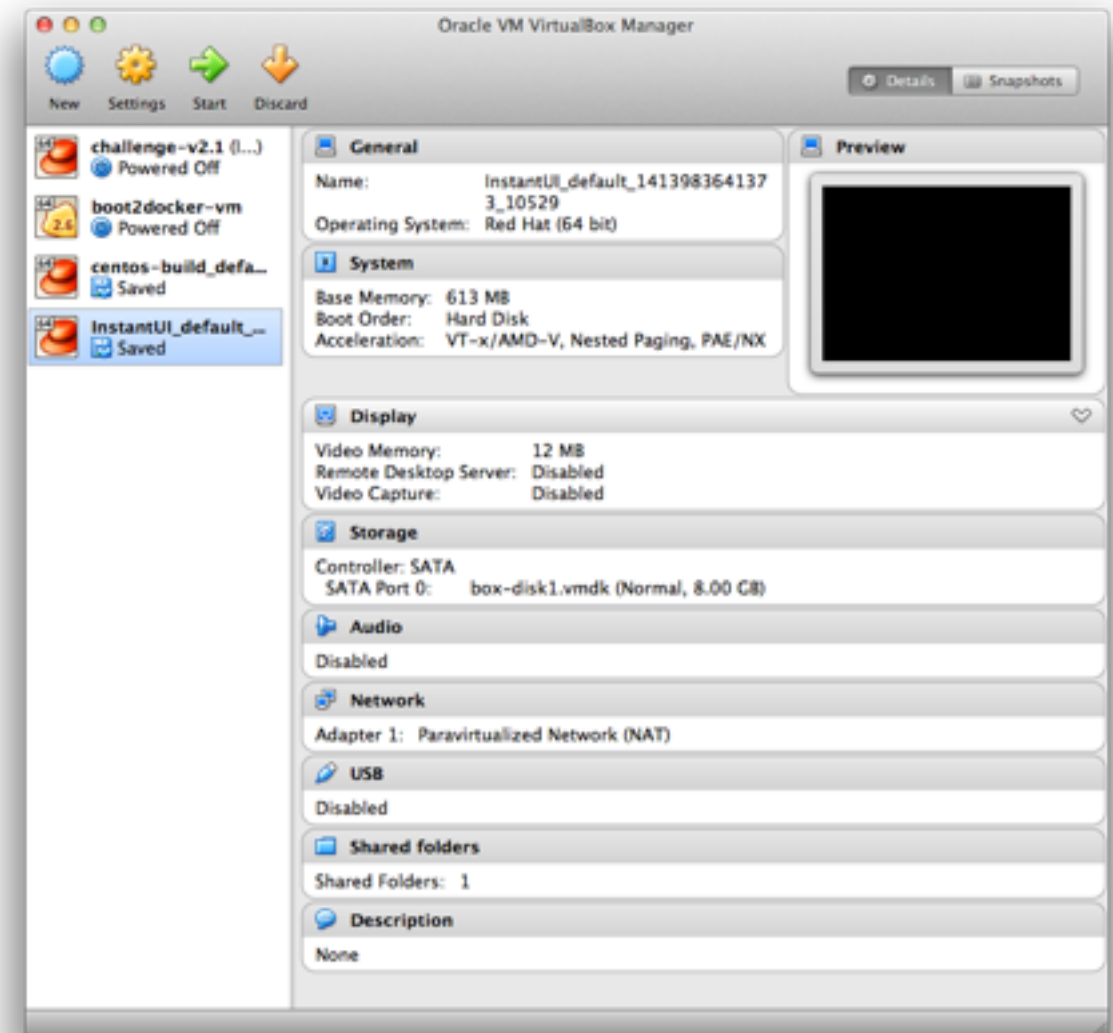
- Because Linux is the kernel you need other pieces of software to make an OS.
- Since most/all of it is open source many people have taken Linux and created different variants, called Distributions.
- There are two main core Distributions - Debian and Fedora (created by RedHat).
- Many other distros exists but in general many of these are “Forks” of either Debian or Fedora.



distrowatch.com

Installing Linux in a VM

- Simplest way to check out Linux.
- VirtualBox is free and available for most platforms (Linux, Windows and OS X).
- Download an ISO of your favourite Distribution and point VirtualBox at it.
- Be sure to give the VM enough memory and enable graphic acceleration for best performance.
- You can also install “guest additions” to allow you to mount directories from your host to the VM.

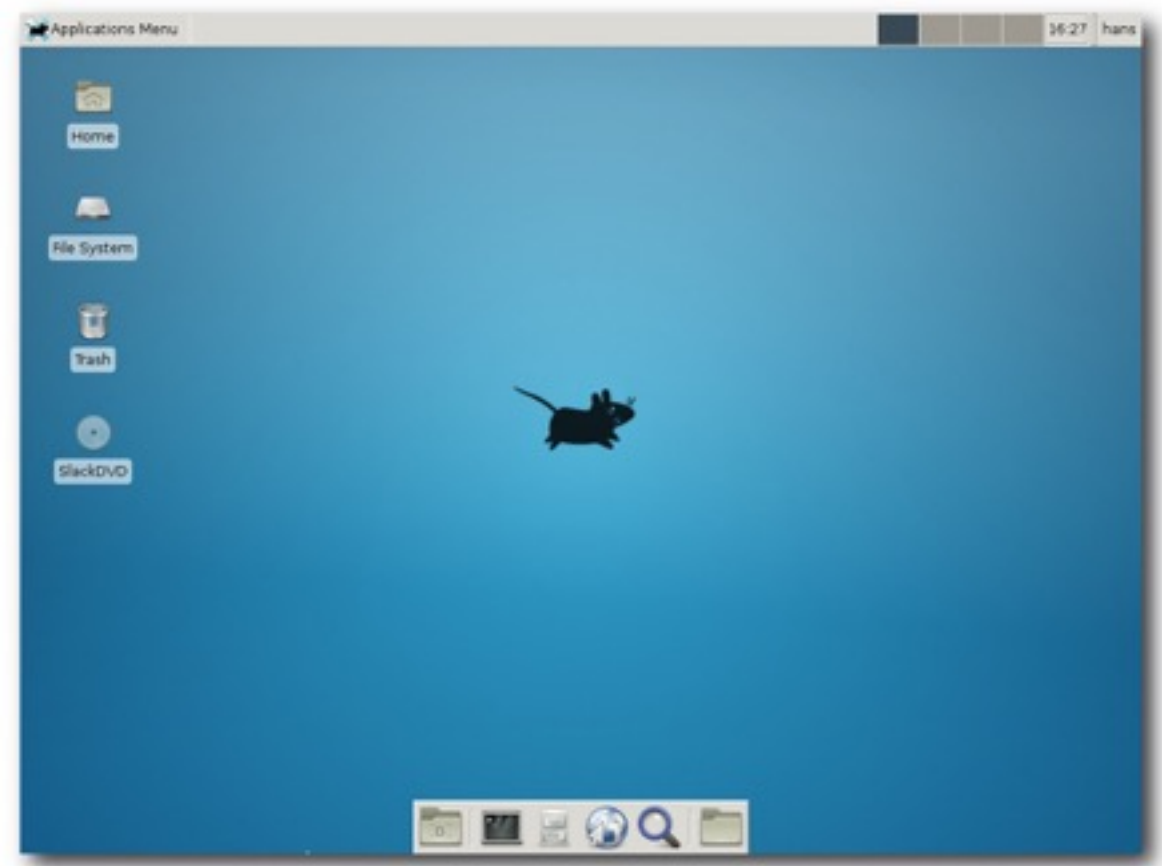


What you'll be using for this course

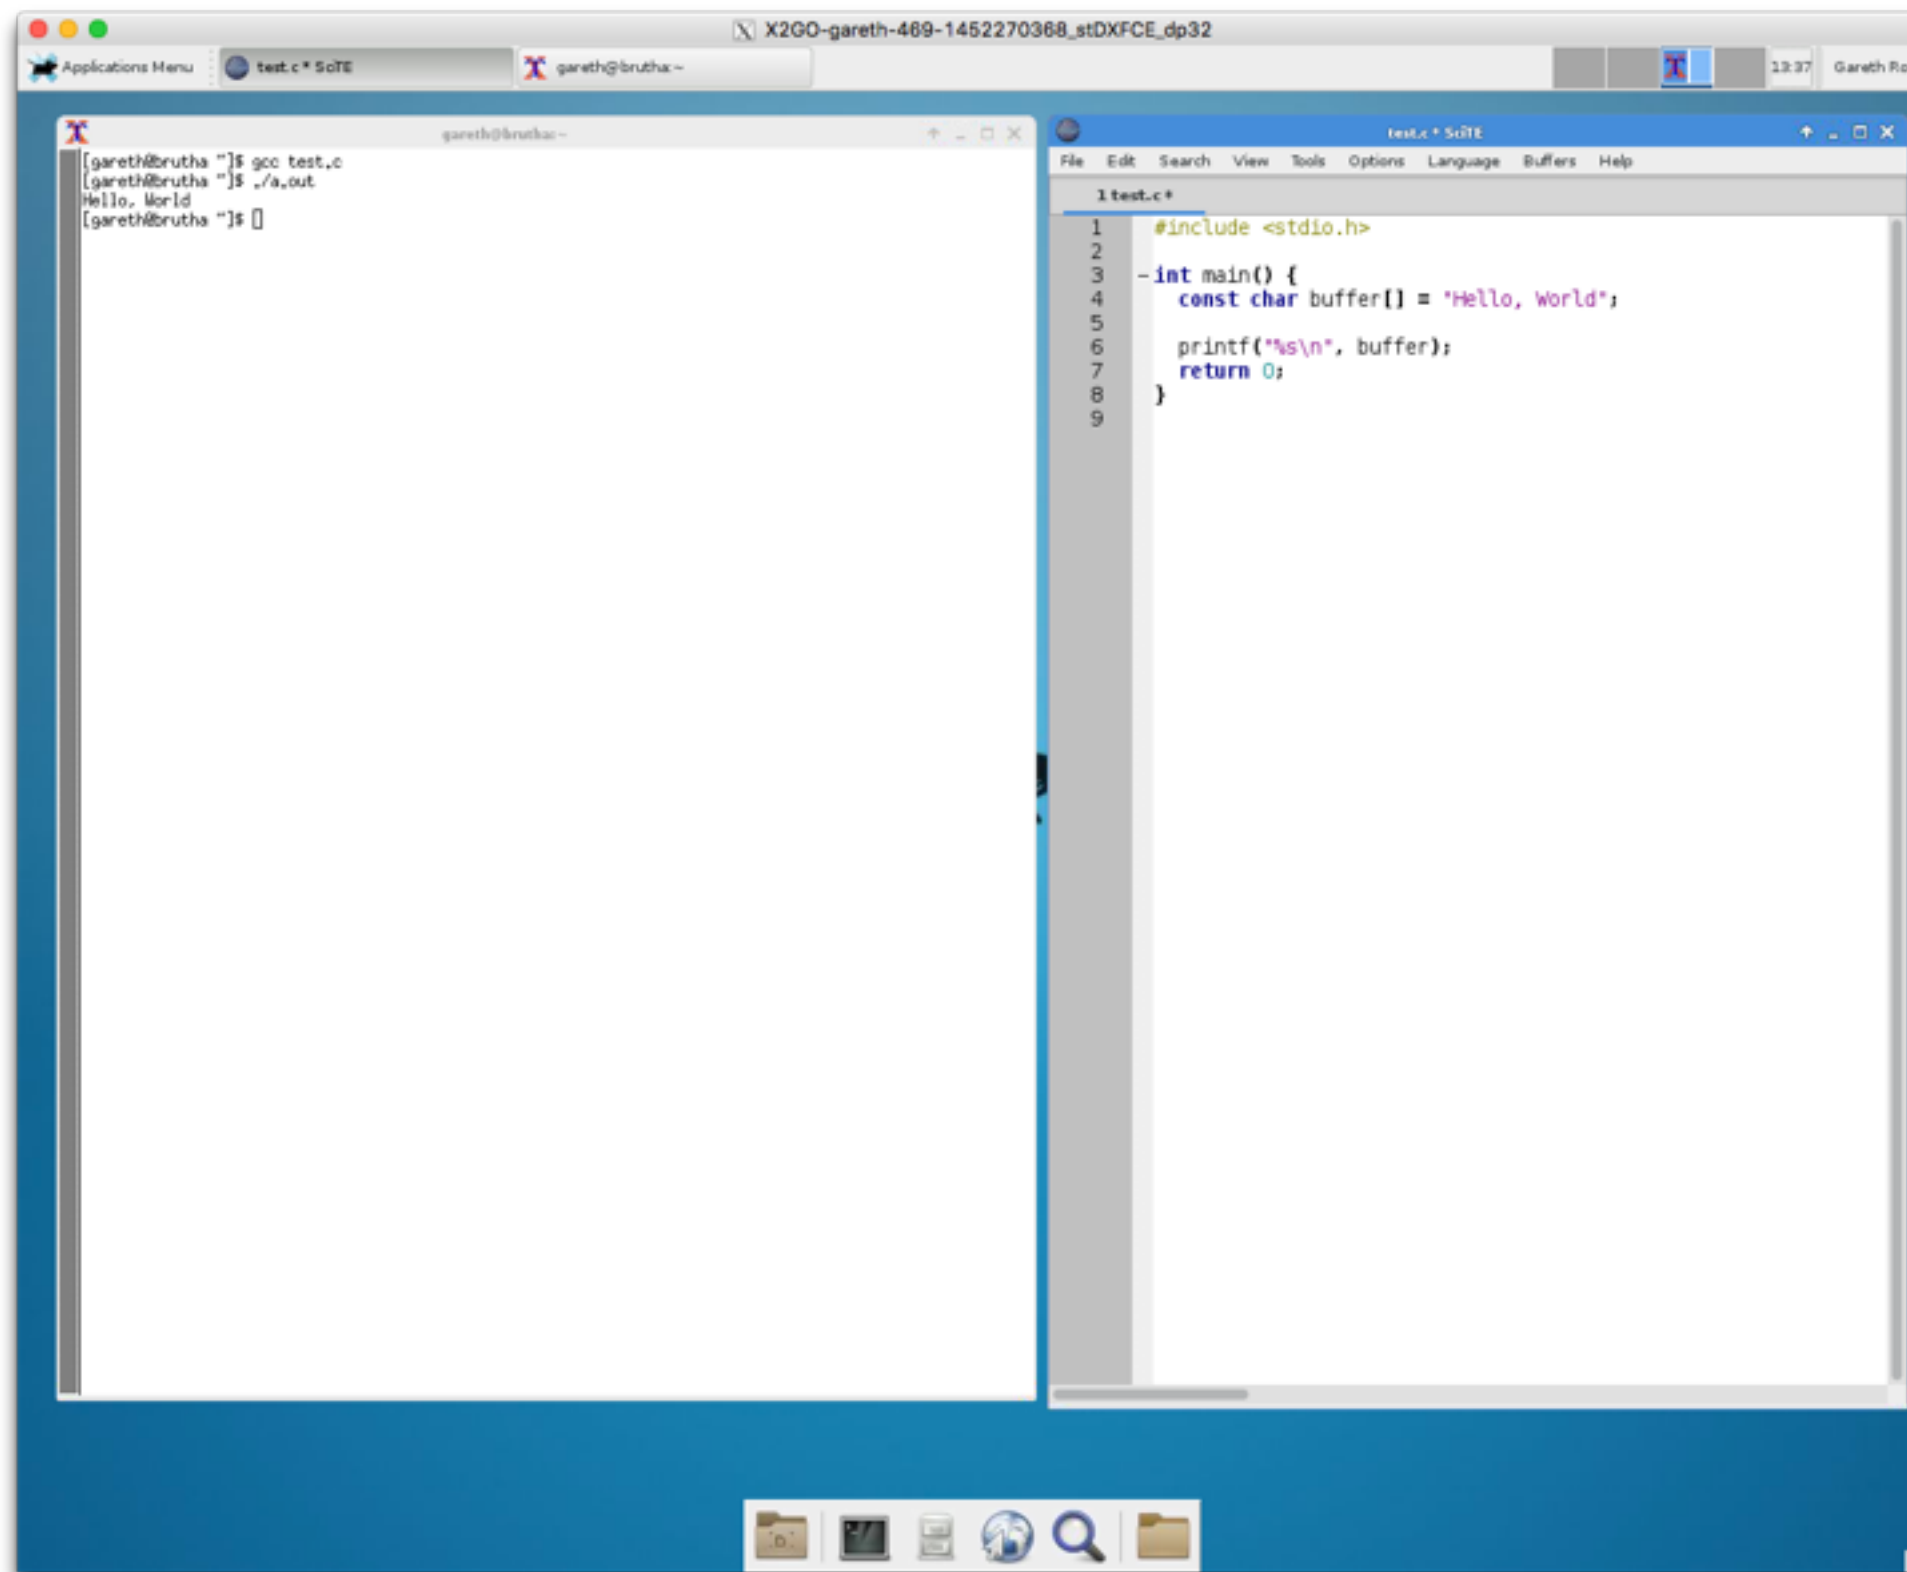
- For this course you'll be given access to remote linux server.
- You'll be using a distribution known as CentOS.
- This is a fork of Redhat's RHEL, popular in business and academia.
- You'll be able to log into the remote server and get a XFCE desktop via a piece of software called X2GO
- Instructions are up on Moodle2



CentOS



Accessing brutha.physics.gla.ac.uk

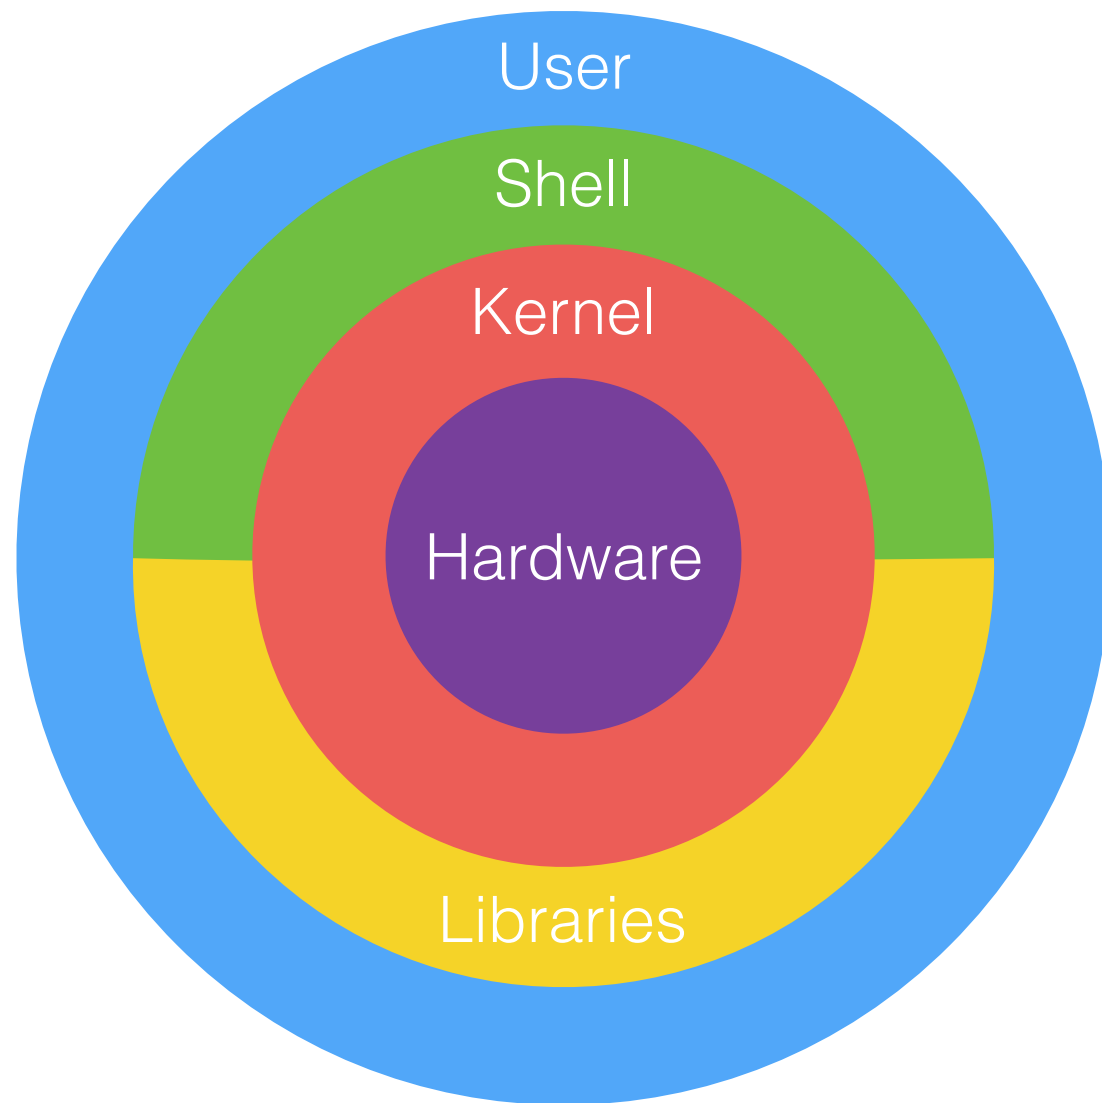


```
[gareth@brutha ~]$ gcc test.c
[gareth@brutha ~]$ ./a.out
Hello, World
[gareth@brutha ~]$
```

```
1 #include <stdio.h>
2
3 int main() {
4     const char buffer[] = "Hello, World";
5
6     printf("%s\n", buffer);
7     return 0;
8 }
9
```

- A brief history lesson
- Why Linux?
- The parts of an Operating System
- The Command Line

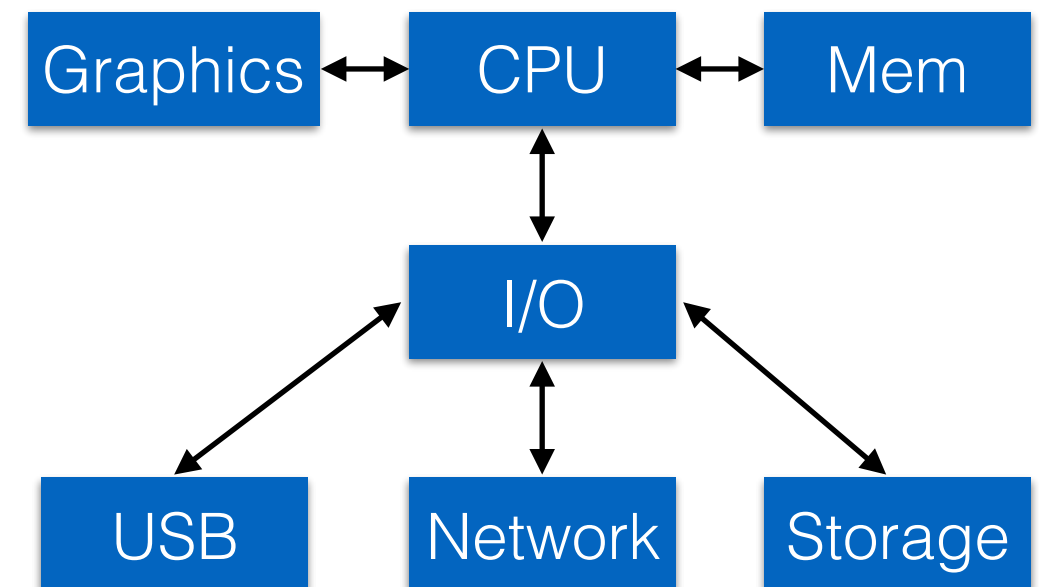
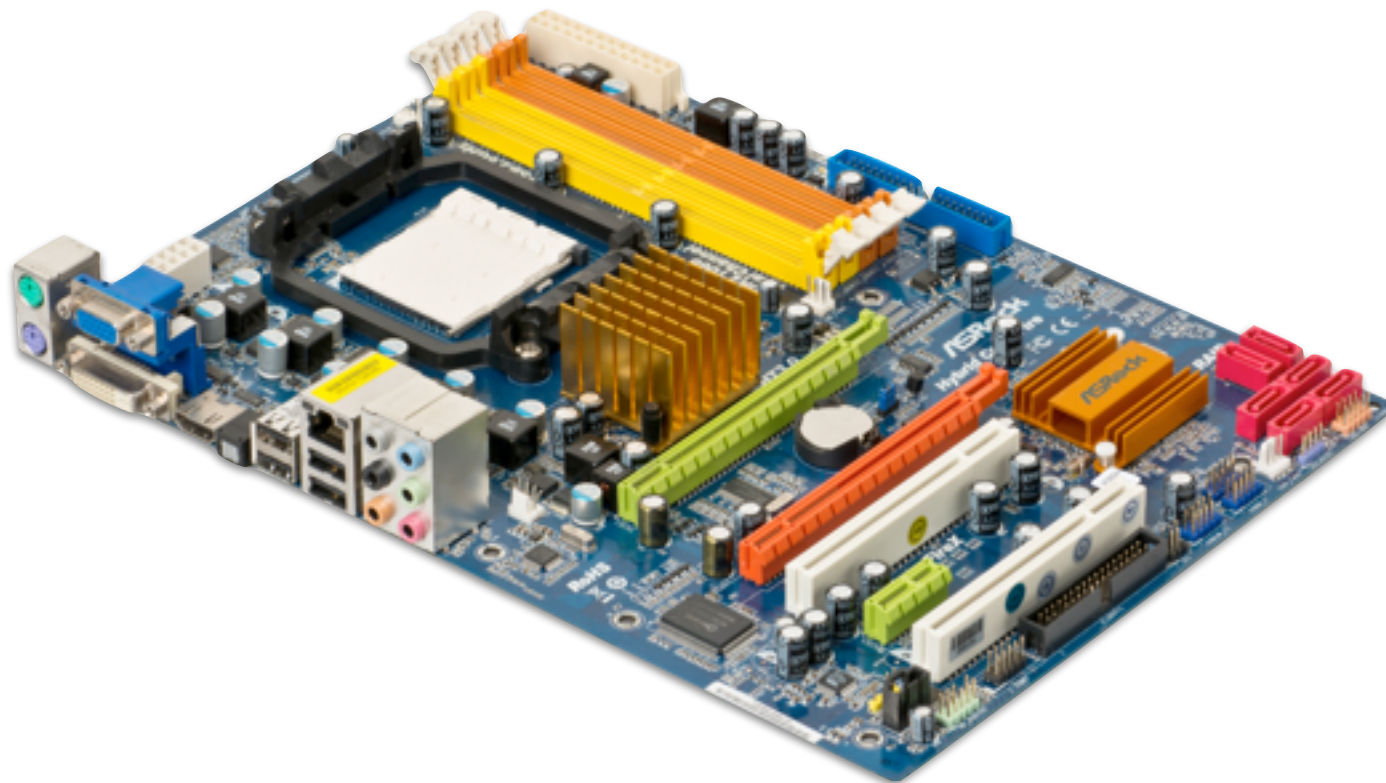
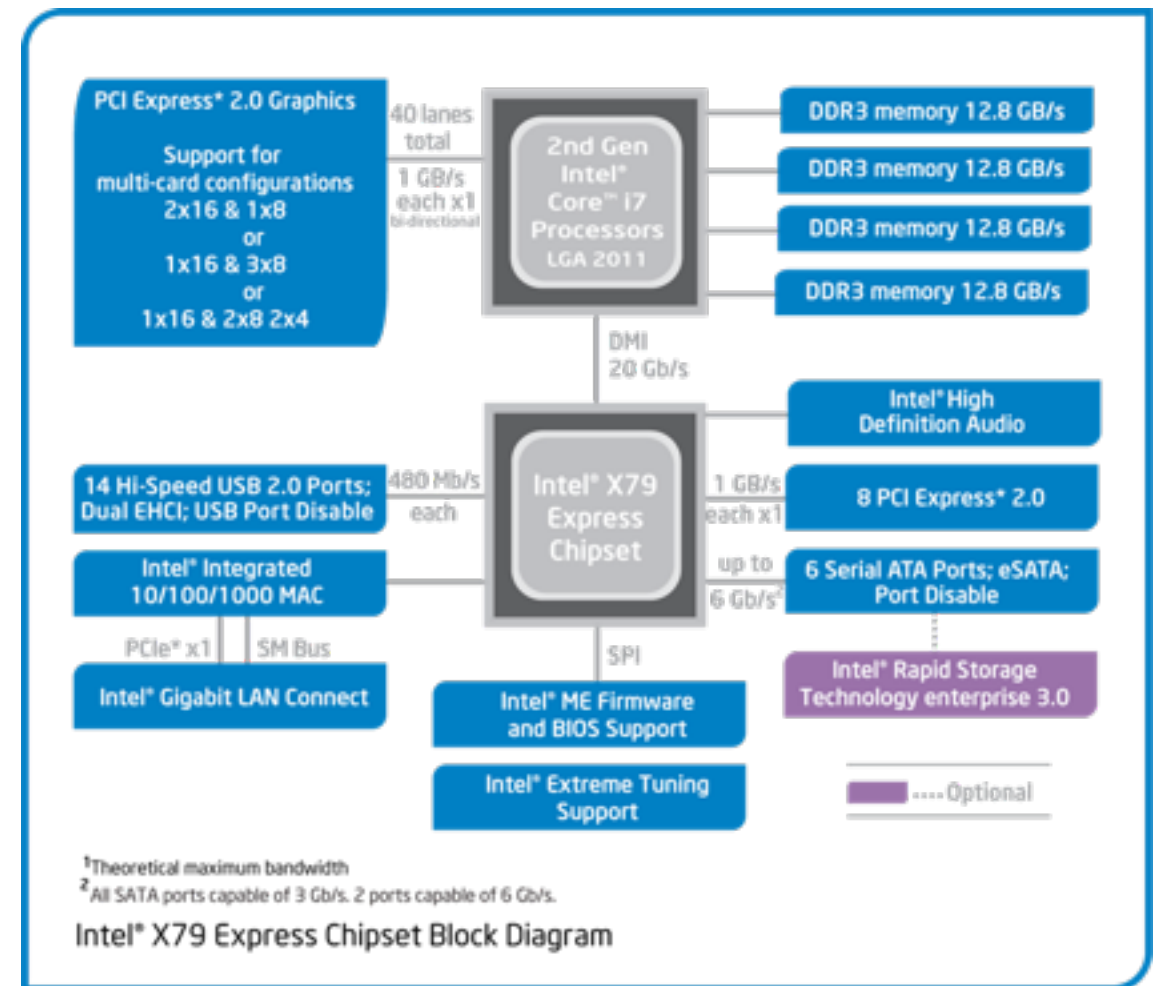
Linux Architecture



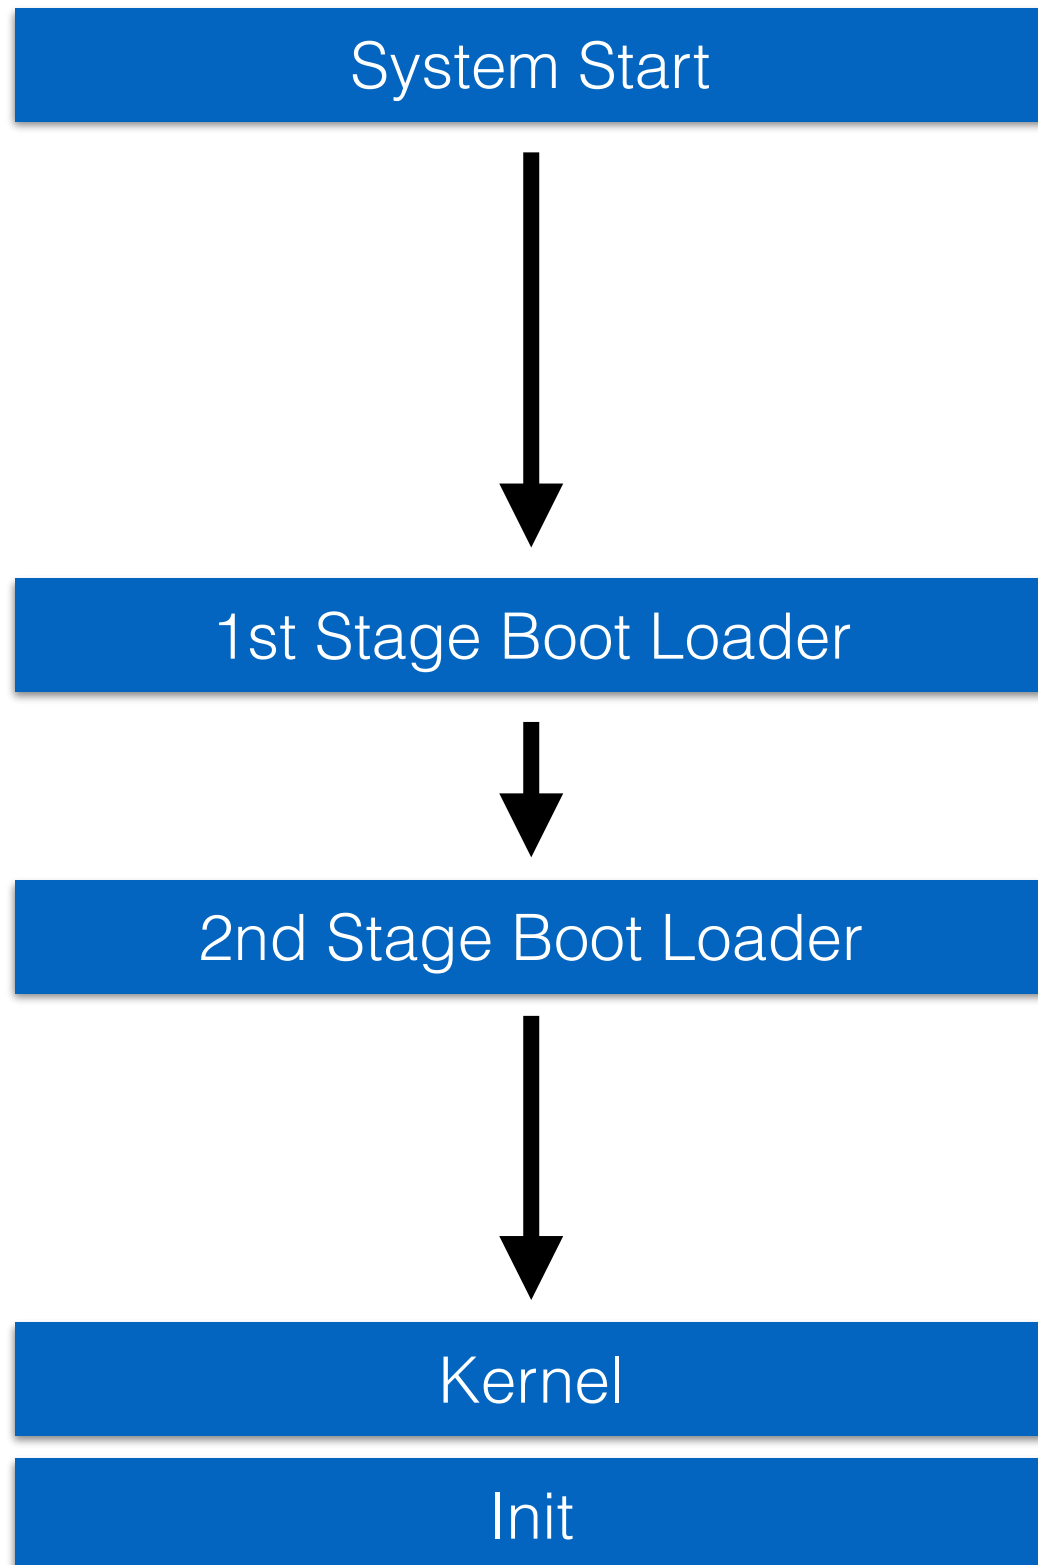
- The Kernel is the core part of Linux, it includes all the drivers need to interact with the underlying hardware.
- The Shell allows interactive access to system resource through a set of builtin commands or by running commands found on the filesystem
- Additionally access to system resources can be via system libraries that expose access to the underlying Kernel (c.f. glibc)
- User applications will usually be started by the shell, or by the init process at system startup.

Hardware

Component	Size	Purpose
CPU	GHz	Numeric / Logic / Shift / Vector Operations
Memory	GB	Running Software and Data storage.
Network	Gbit	Data transmission, TCP/IP packet decoding.
Storage	TB	Data storage preserved without power requirements
Graphics	MOp/s	VGA Display, 2D/3D Acceleration, Floating point ops.



Boot Process



Power On Self Test (POST)

- Verify BIOS code
- Verify CPU and Registers
- Find and Verify Memory
- Initialise the BIOS
- Discover and Initialise Hardware

Master Boot Record

- 512 bytes located on first sector
- Small program that loads 2nd Stage


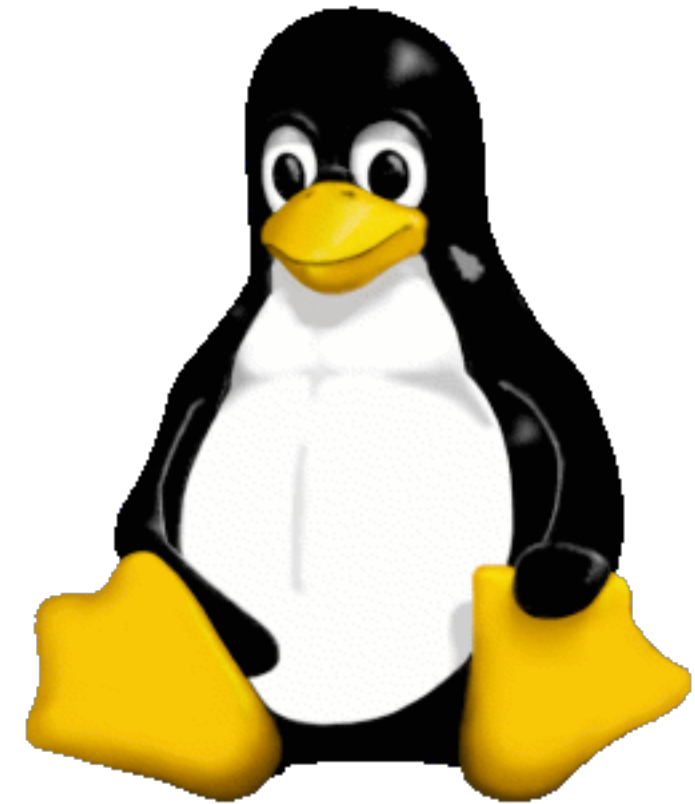
Loads the Kernel

- Loads required components into Mem.
- Enters Protected mode
- Transfers control to the Kernel

Kernel starts and brings up system services. Once devices have been initialised and important subsystems started, starts init services.

Kernel

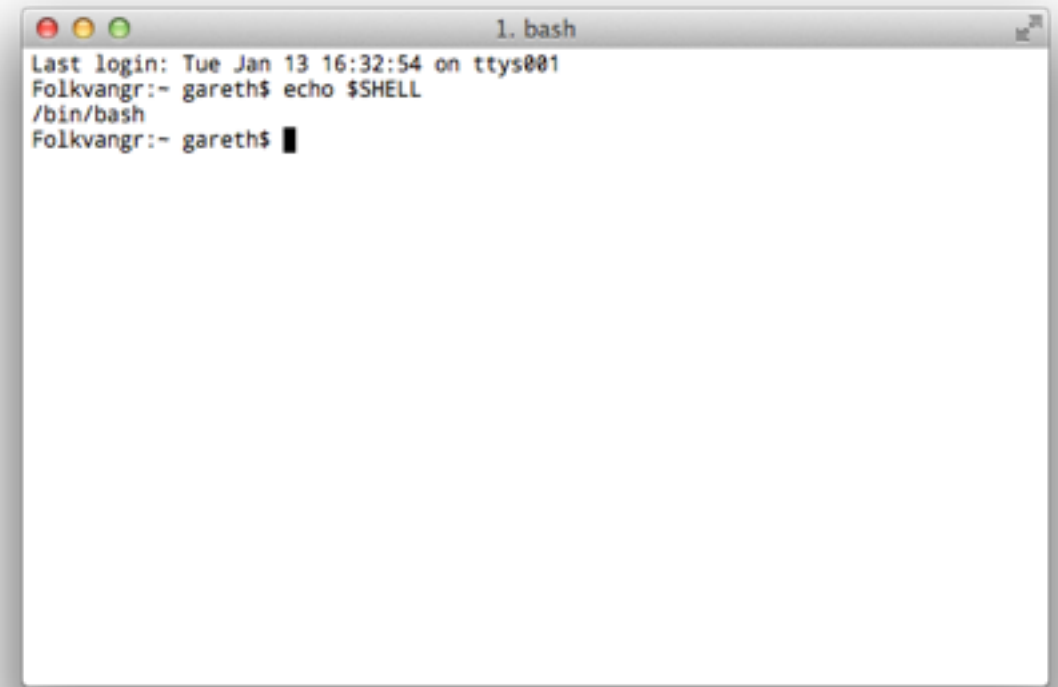
- The Kernel looks after most interactions with Hardware and all of the process running on the system
- It provides:
 - Process scheduler
 - Inter-Process Communication
 - Memory Management
 - A Virtual Filesystem
 - Networking
 - Device Mapper
 - Sound Subsystem
 - etc...



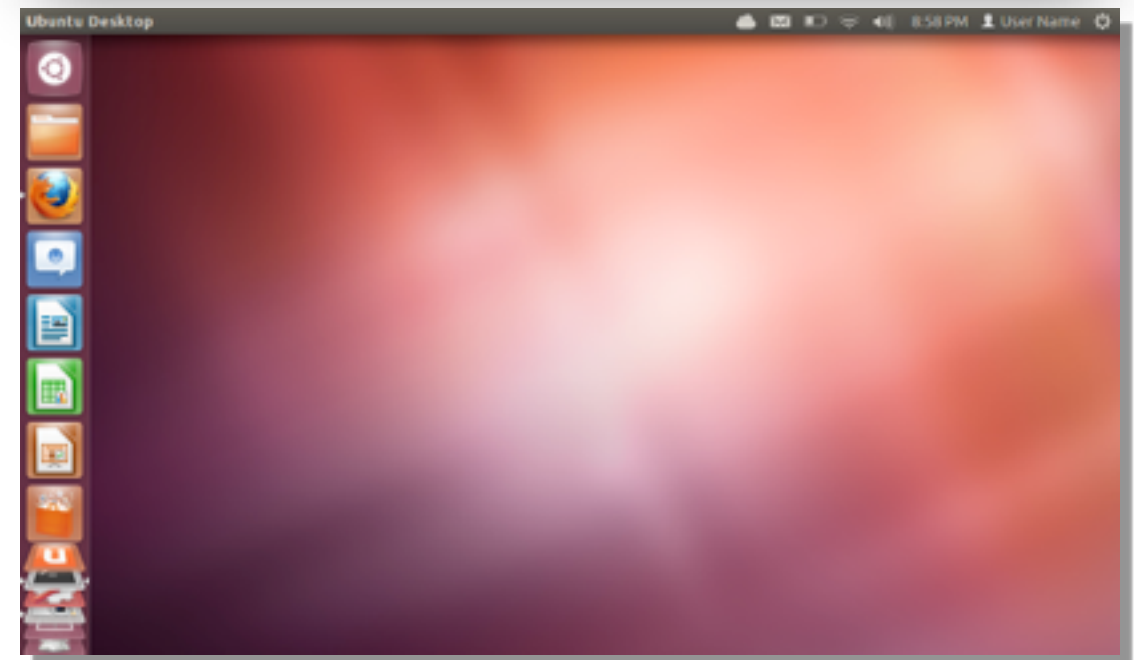
```
isapnp: No Plug & Play device found
Linux NET4.0 for Linux 2.4
Based upon Swansea University Computer Society NET3.039
initializing RT netlink socket
Starting kswapd
VFS: Disk quotas vfstab 6.5.1
vesafb: framebuffer at 0xc0000000, mapped to 0xc2004000, size 51200k
vesafb: mode is 800x600x8, linelength=800, pages=3
vesafb: protected mode interface info at 45f3:1f5f
vesafb: scrolling: redraw
Console: switching to colour frame buffer device 100x37
fb0: VESA VGA frame buffer device
pty: 256 Unix98 pty's configured
Uniform Multi-Platform E-IDE driver Revision: 6.31
ide: Assuming 50MHz system bus speed for PIO modes; override with idebus=xx
hda: Generic 1234, ATAPI CD-ROM drive
ide0 at 0x1f0-0x1f7,0x3f6 on irq 14
hda: ATAPI 4X CD-ROM drive, 512kB Cache
Uniform CD-ROM driver Revision: 3.12
FPC 0 is an 8272A
RAMDISK driver initialized: 16 RAM disks of 4096K size 1024 blocksize
Cromx Ltd. Synchronous PPP and CISCO HDLC (c) 1994
Linux port (c) 1998 Building Number Three Ltd & Jan "Feng" Keszrak.
SCSI subsystem driver Revision: 1.00
scsi0 : SCSI host adapter emulation for IDE ATAPI devices
NET4: Linux TCP/IP 1.0 for NET4.0
IP Protocols: ICMP, UDP, TCP, IGMP
IP: routing cache hash table of 512 buckets, 4Kbytes
TCP: Hash tables configured (established 2048 bind 2048)
NET4: Unix domain sockets 1.0/SMP for Linux NET4.0.
RAMDISK: Compressed image found at block 0
```

Shell

- In Linux the shell handles interaction between Users and the system.
- It can be via a text interface known as the command line as well as through a graphical interface, for instance the “gnome-shell”
- It allows other processes/applications to be started such as editors, compilers, office suites.
- We’ll cover the command line in more detail later in this set of slides.

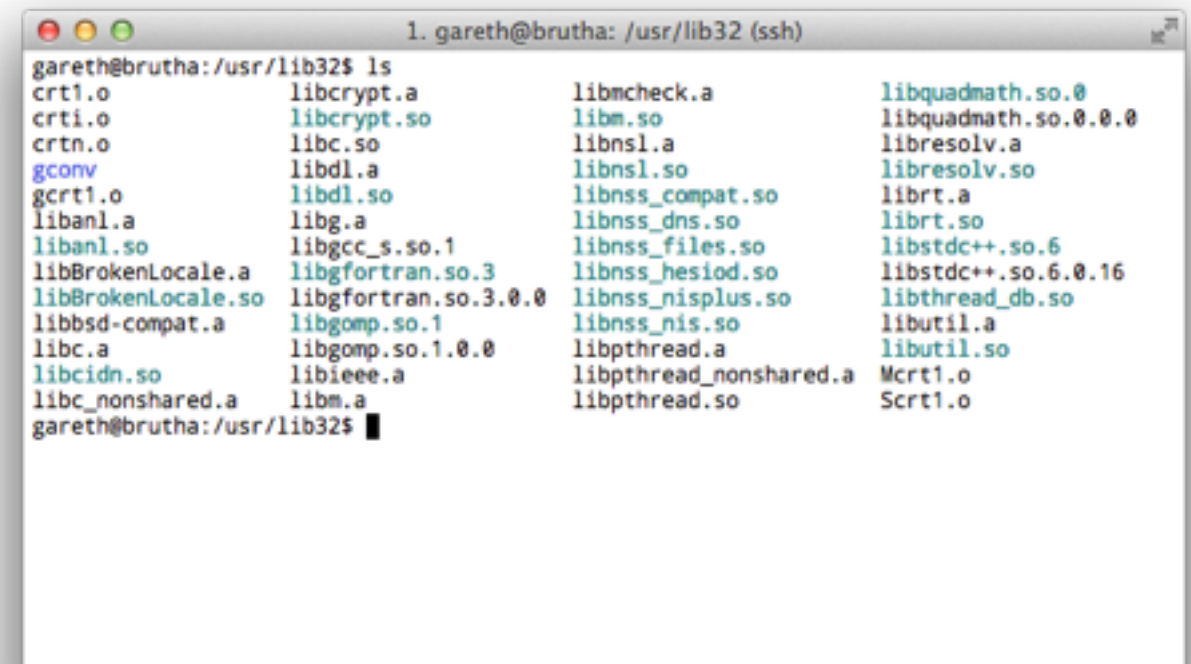
A terminal window titled "1. bash" showing a login session. The text inside the terminal is: "Last login: Tue Jan 13 16:32:54 on ttys001", "Folkvangr:~ gareth\$ echo \$SHELL", "/bin/bash", and "Folkvangr:~ gareth\$".

```
1. bash
Last login: Tue Jan 13 16:32:54 on ttys001
Folkvangr:~ gareth$ echo $SHELL
/bin/bash
Folkvangr:~ gareth$
```



System Libraries

- Running programs also need to interact with system resources.
- This is carried out via system libraries.
- There are a wide range of these most of which are beyond the scope of this course.
- You will however use the standard C libraries as part of the C programming component of this course.



```
1. gareth@brutha: /usr/lib32 (ssh)
gareth@brutha:/usr/lib32$ ls
crt1.o      libcrypt.a  libmcheck.a  libquadmath.so.0
crti.o      libcrypt.so  libm.so       libquadmath.so.0.0.0
crtn.o      libc.so     libnsl.a      libresolv.a
gconv       libdl.a     libnsl.so     libresolv.so
gcrt1.o     libdl.so    libnss_compat.so  librt.a
libanl.a    libg.a      libnss_dns.so  librt.so
libanl.so   libgcc_s.so.1  libnss_files.so  libstdc++.so.6
libBrokenLocale.a  libgfortran.so.3  libnss_hesiod.so  libstdc++.so.6.0.16
libBrokenLocale.so  libgfortran.so.3.0.0  libnss_nisplus.so  libthread_db.so
libbsd-compat.a  libgomp.so.1      libnss_nis.so     libutil.a
libc.a       libgomp.so.1.0.0  libnss_nisplus.so  libutil.so
libcidn.so   libieee.a      libpthread.a     libcrt1.o
libc_nonshared.a  libm.a         libpthread_nonshared.a  Scrt1.o
libpthread.so
```


Userland

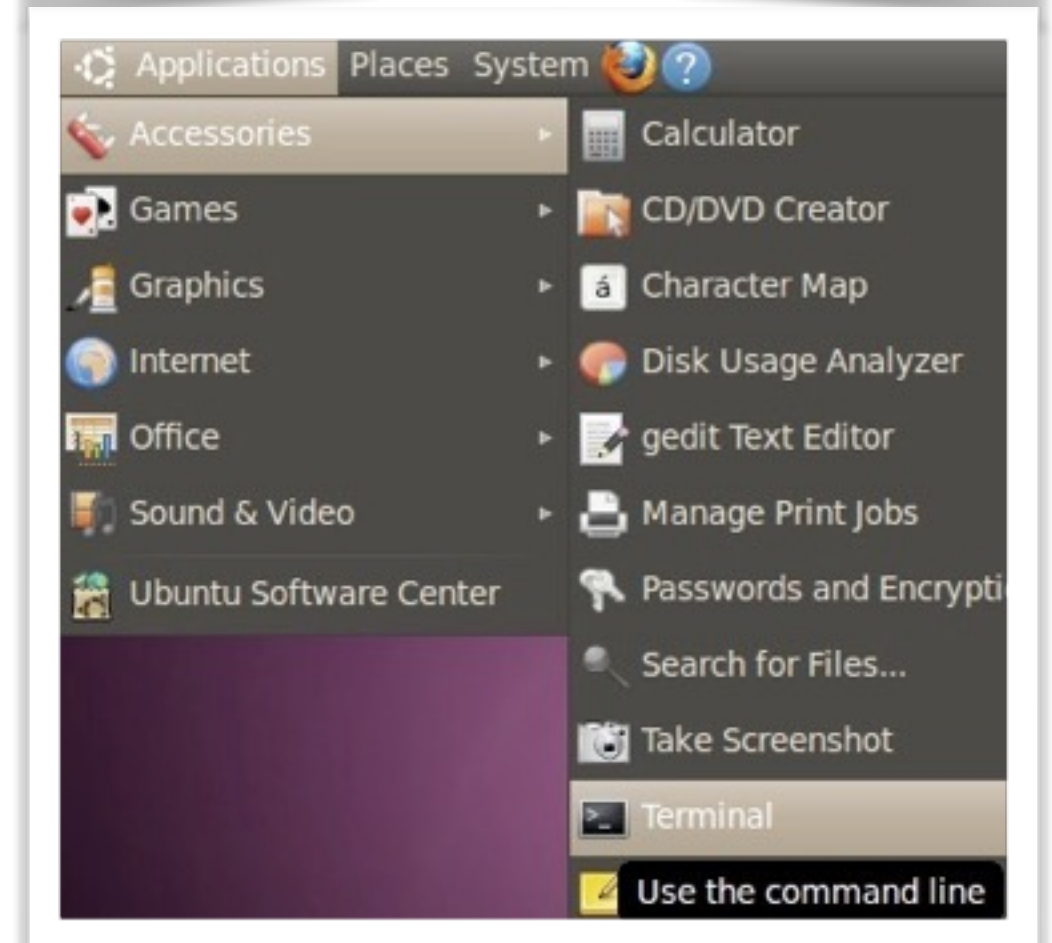
- User processes are the final level of the Operating System.
- This is likely the most familiar and includes programs such as web browsers, office suites, editors etc.
- Most user orientated distributions ship with software that allows you to do most common tasks.



- A brief history lesson
- Why Linux?
- The parts of an Operating System
- The Command Line

The Command Line

- There are a variety of different shells available: Sh, BASH, CSh, TCSH, ZSh...
- C-style shells: C-like syntax for scripting
- Bash-style shells: slightly more consistent, and more widely adopted in the the research-computing community.
- We'll use Bash in this course, but feel free to experiment!
- To open the terminal, click the ubuntu button and type "terminal"!

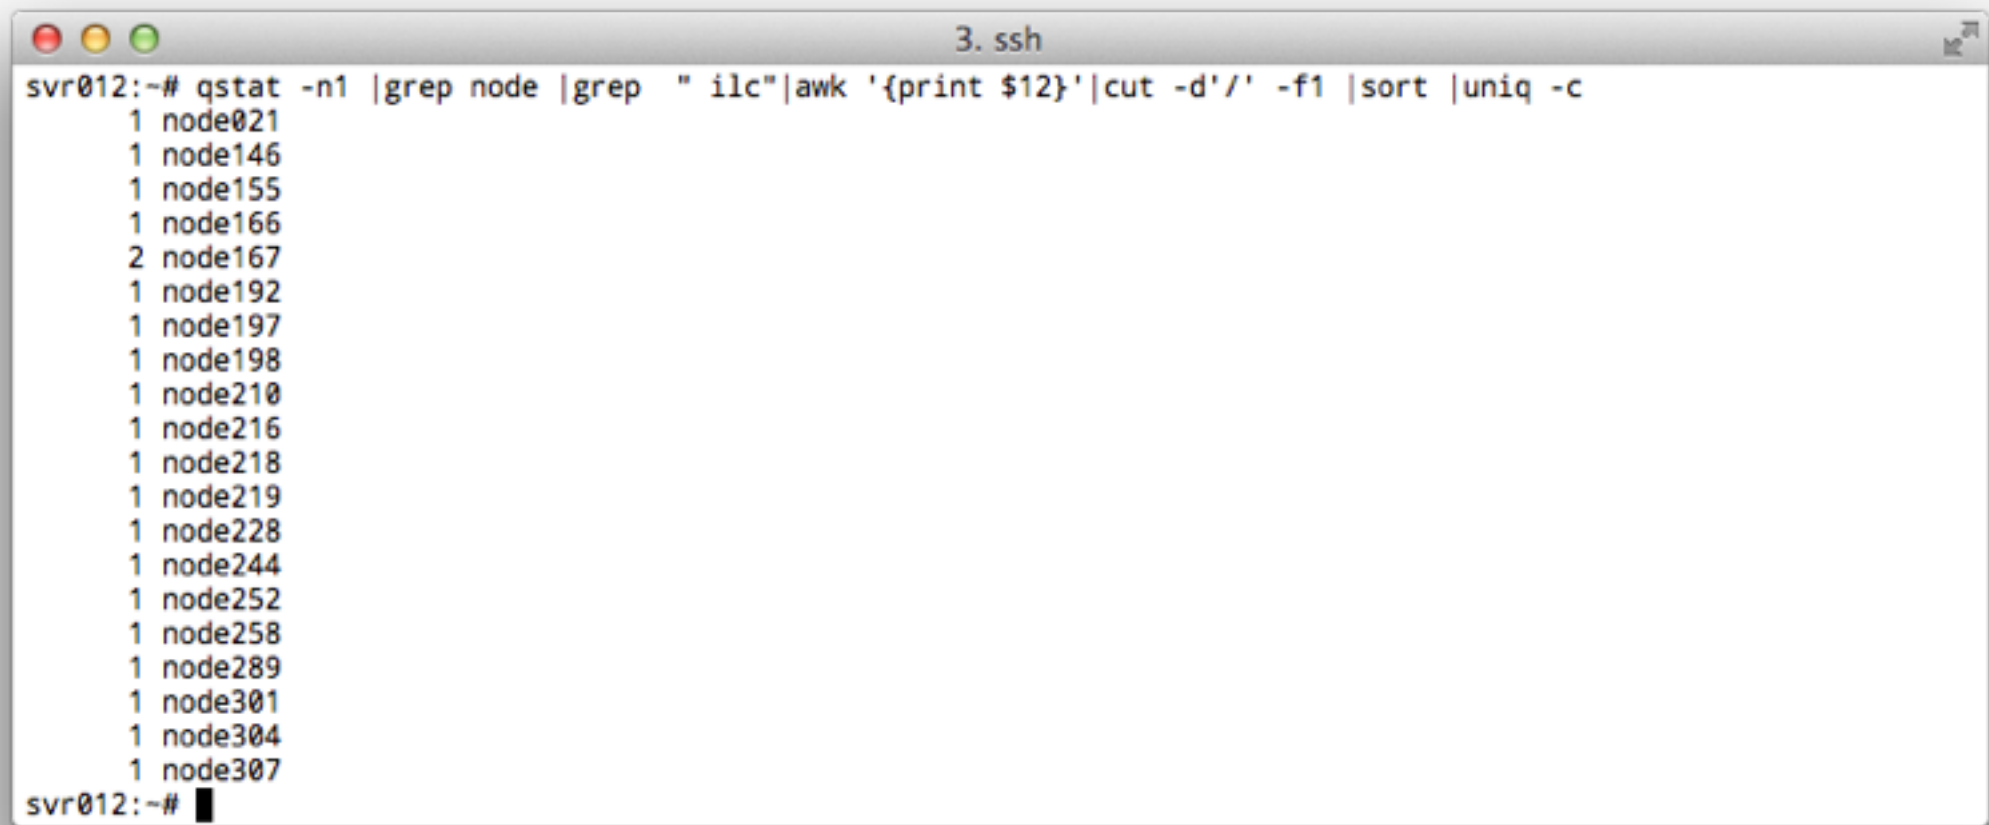


The Command Line

- The command line lets us input command at the shell prompt and receive output in a text form.
- Commands are of the form:

prompt# <command> <flags> <arguments>

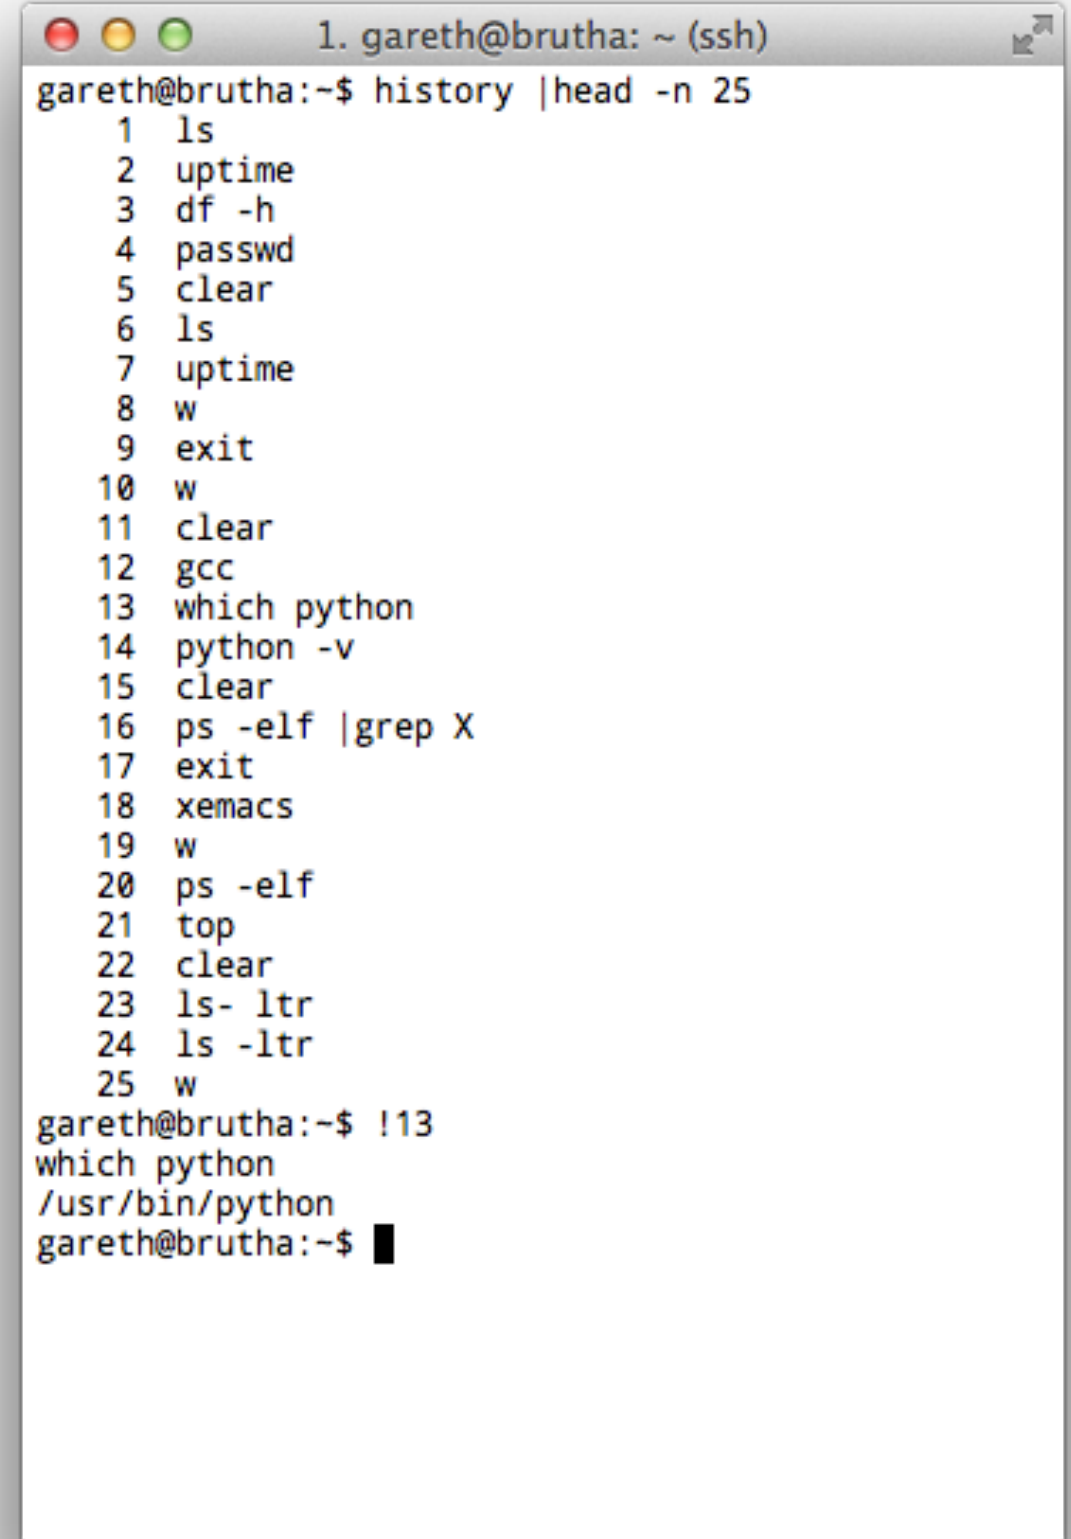
prompt# du -h --max-depth=1
- Commands can be chained together to build up complex interactions and obtain the results we want.



```
3. ssh
svr012:~# qstat -n1 |grep node |grep " ilc"|awk '{print $12}'|cut -d'/' -f1 |sort |uniq -c
  1 node021
  1 node146
  1 node155
  1 node166
  2 node167
  1 node192
  1 node197
  1 node198
  1 node210
  1 node216
  1 node218
  1 node219
  1 node228
  1 node244
  1 node252
  1 node258
  1 node289
  1 node301
  1 node304
  1 node307
svr012:~#
```

Useful Commands

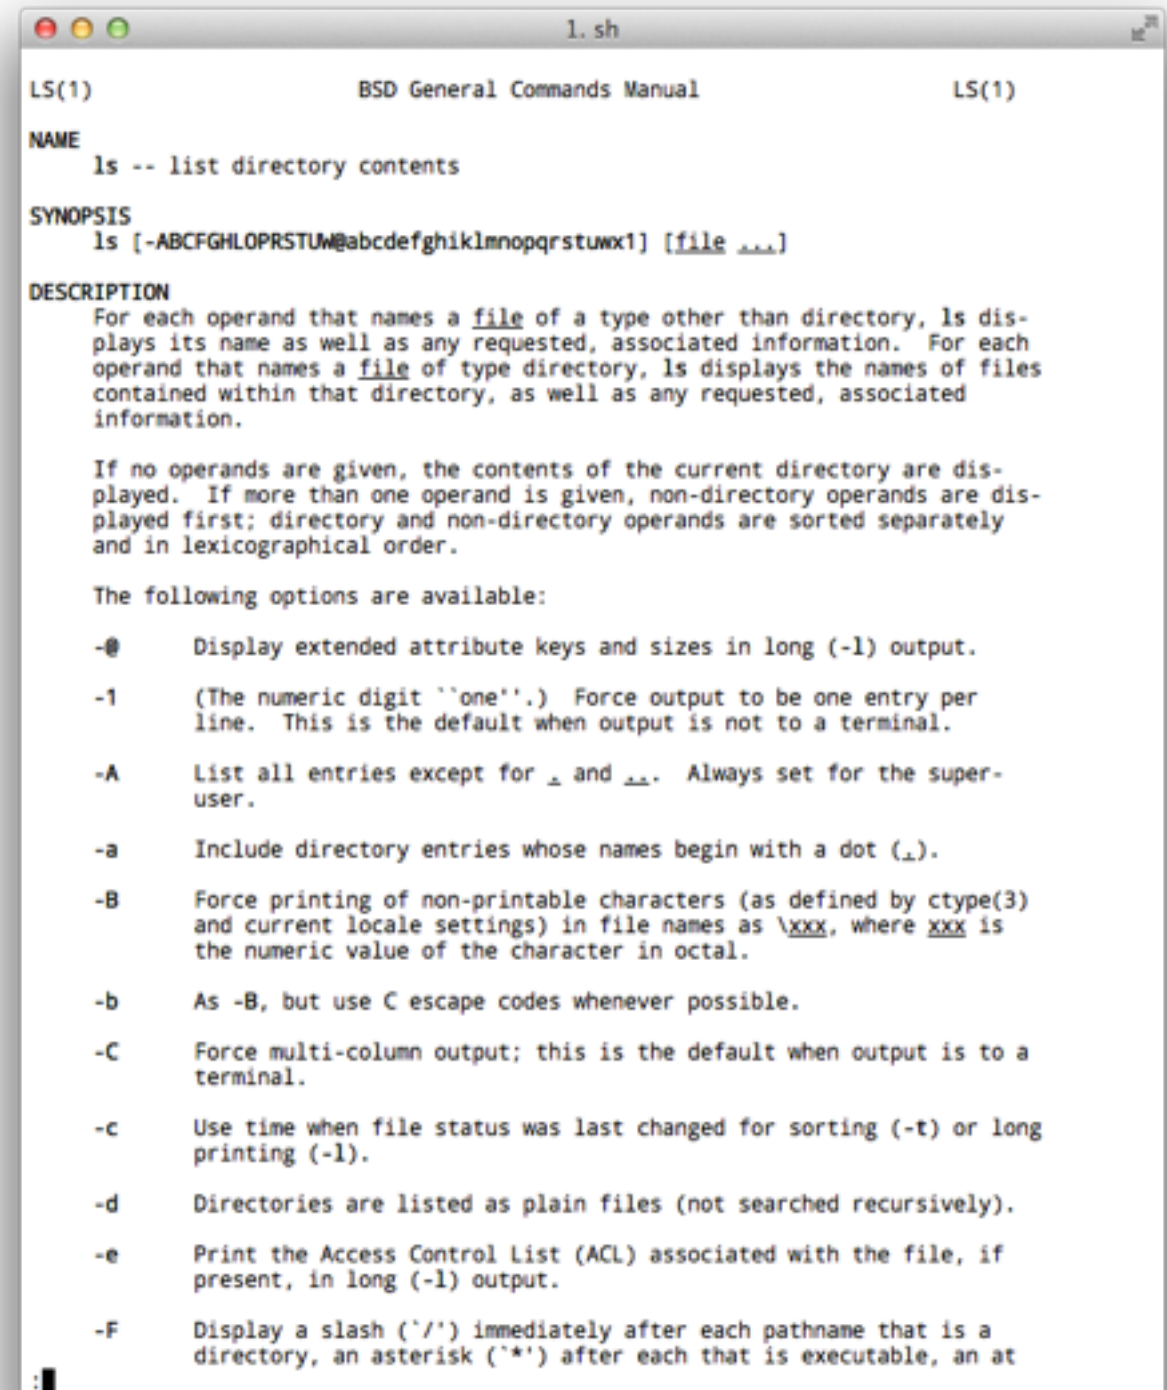
- **history** - gives all commands you've entered
- **!!** - runs the last command
- **!**<number>**** runs the command at <number> i.e. **!484** runs command 484
- **up/down arrow** cycles through previous commands
- **tab** auto completes names



```
1. gareth@brutha: ~ (ssh)
gareth@brutha:~$ history |head -n 25
 1  ls
 2  uptime
 3  df -h
 4  passwd
 5  clear
 6  ls
 7  uptime
 8  w
 9  exit
10  w
11  clear
12  gcc
13  which python
14  python -v
15  clear
16  ps -elf |grep X
17  exit
18  xemacs
19  w
20  ps -elf
21  top
22  clear
23  ls- ltr
24  ls -ltr
25  w
gareth@brutha:~$ !13
which python
/usr/bin/python
gareth@brutha:~$
```

Useful Commands (2)

- **man** will give you the manual page for a particular command.
 - hitting the **spacebar** takes you to the next page and **q** quits
- **ctrl-r** allows you to search through previous commands
- **clear** empties the screen (useful for clearing your head sometimes).



```
1. sh
LS(1)                                BSD General Commands Manual                                LS(1)
NAME
    ls -- list directory contents
SYNOPSIS
    ls [-ABCFGHLOPRSTUW@abcdefghiklmnopqrstuwx1] [file ...]
DESCRIPTION
    For each operand that names a file of a type other than directory, ls displays its name as well as any requested, associated information. For each operand that names a file of type directory, ls displays the names of files contained within that directory, as well as any requested, associated information.
    If no operands are given, the contents of the current directory are displayed. If more than one operand is given, non-directory operands are displayed first; directory and non-directory operands are sorted separately and in lexicographical order.
    The following options are available:
    -@      Display extended attribute keys and sizes in long (-l) output.
    -1      (The numeric digit 'one'.) Force output to be one entry per line. This is the default when output is not to a terminal.
    -A      List all entries except for . and ... Always set for the super-user.
    -a      Include directory entries whose names begin with a dot (.).
    -B      Force printing of non-printable characters (as defined by ctype(3) and current locale settings) in file names as \xxx, where xxx is the numeric value of the character in octal.
    -b      As -B, but use C escape codes whenever possible.
    -C      Force multi-column output; this is the default when output is to a terminal.
    -c      Use time when file status was last changed for sorting (-t) or long printing (-l).
    -d      Directories are listed as plain files (not searched recursively).
    -e      Print the Access Control List (ACL) associated with the file, if present, in long (-l) output.
    -F      Display a slash ('/') immediately after each pathname that is a directory, an asterisk ('*') after each that is executable, an at
```

- A brief history lesson
- Why Linux?
- The parts of an Operating System
- The Command Line